

A Comprehensive Literature Review with Self-Reflection

Literature Review

October 7, 2025

Abstract

This literature review provides a comprehensive analysis of recent research in the field. The review synthesizes findings from 146 research papers, identifying key themes, methodological approaches, and future research directions.

Contents

1 Introduction

1.1 Historical Trajectory and Early Promise of RL in NLP

The application of reinforcement learning (RL) principles to natural language processing (NLP) emerged from a fundamental challenge inherent in traditional supervised learning: the limitations of maximum likelihood estimation (MLE) and its inability to directly optimize non-differentiable, task-specific evaluation metrics. Supervised models, while effective for token-level prediction, often suffered from "exposure bias" during sequence generation. This phenomenon occurs when a model, trained on ground-truth sequences, is forced to generate tokens conditioned on its *own* potentially erroneous previous outputs during inference, leading to a compounding of errors and a drift from the desired output distribution ?. Such models frequently produced generic, repetitive, or grammatically plausible but semantically uninspired outputs, failing to capture the nuanced, holistic quality desired in human language.

This critical limitation motivated a paradigm shift: viewing many NLP tasks, particularly those involving structured outputs or generative sequences, as sequential decision-making processes. In this framework, an "agent" learns to construct an output by making a series of choices (e.g., selecting the next word or structural element), receiving delayed rewards based on the quality of the complete structure or generated sequence, rather than just local token-level accuracy. This ability to learn from delayed, global rewards offered a distinct advantage over MLE, which optimizes local, independent probabilities and struggles to account for long-term dependencies and overall output quality.

Early conceptualizations of RL in NLP primarily focused on structured prediction tasks such as parsing, semantic role labeling, and sequence tagging. Here, the objective was to optimize for global structural correctness, often measured by non-differentiable metrics like F1-score or exact match, rather than just individual token or constituent accuracy. A crucial intellectual bridge in this era was the "learning to search" (L2S) framework, exemplified by approaches like SEARN ? and DAgger ?. These methods elegantly reframed complex structured prediction as a series of local decisions made by a

learned policy, often leveraging imitation learning to guide the search process. While not pure RL in the contemporary sense, L2S provided a robust conceptual and algorithmic foundation for sequential decision-making in NLP, demonstrating how a policy could learn to navigate a complex output space to optimize for a global objective, thereby laying the groundwork for more direct RL applications. This perspective highlighted the sequential nature of constructing linguistic outputs and the need for global optimization beyond local token-level accuracy, a theme that would become central to later Deep Reinforcement Learning (DRL) applications.

As neural networks and deep learning gained prominence in the mid-2010s, the principles of RL were extended to deep models, giving rise to DRL for direct language generation. This era, preceding the widespread adoption of large language models (LLMs), saw initial demonstrations of DRL’s potential to address the limitations of MLE in generative tasks such as dialogue generation, abstractive summarization, and neural machine translation. The core idea was to train neural sequence-to-sequence models to directly optimize for human-aligned, non-differentiable metrics (e.g., ROUGE for summarization, BLEU for translation, or task-completion rates for dialogue) by treating the generation process as a Markov Decision Process. Policy gradient methods, such as REINFORCE, were adapted to train recurrent neural networks to generate sequences by maximizing a reward signal derived from these holistic metrics ?. This allowed models to learn from the consequences of their entire generated sequence, rather than just the likelihood of individual tokens, leading to outputs that were often more fluent, coherent, and aligned with overall task objectives than their MLE-trained counterparts.

Despite this early promise, the application of DRL to NLP presented significant challenges that limited its widespread adoption before the LLM era. The high variance inherent in policy gradient methods, especially when operating in the immense discrete action spaces of natural language (i.e., the vocabulary size, often tens of thousands of tokens), frequently led to unstable and inefficient training. Furthermore, the problem of sparse rewards meant that meaningful feedback was often only available at the end of a long sequence, making credit assignment difficult across numerous sequential decisions.

Designing effective, non-brittle reward functions for subjective qualities like fluency, coherence, or relevance proved to be a labor-intensive and often domain-specific engineering endeavor, requiring careful tuning and often proxy metrics that did not perfectly align with human judgment. The inherent sample inefficiency of RL algorithms also posed a significant hurdle, requiring extensive interaction with the environment (or a carefully constructed simulator) to learn an effective policy, a costly process for complex language tasks. These persistent issues—training instability, reward engineering complexity, and sample inefficiency exacerbated by the unique characteristics of language—underscored the need for more robust algorithms and sophisticated feedback mechanisms. These unresolved challenges ultimately set the stage for the next wave of RL innovations in NLP, particularly the development of more stable DRL algorithms and, crucially, the integration of RL with human feedback for aligning large language models.

1.2 Contemporary Motivations: Aligning Large Language Models with Human Intent

The remarkable generative capabilities of Large Language Models (LLMs), primarily achieved through large-scale unsupervised pre-training and subsequent supervised fine-tuning (SFT), often fall short of producing outputs that consistently align with complex human preferences, ethical guidelines, and nuanced instructions ?? . Traditional supervised methods, while effective for learning factual knowledge and grammatical structures, struggle to capture subjective qualities such as helpfulness, harmlessness, coherence, trustworthiness, and intricate instruction following. This limitation stems from their reliance on static, human-annotated datasets, which are inherently limited in scale, prone to biases, and unable to adapt to the dynamic, context-dependent nature of human judgment ?? . This fundamental gap between raw generative power and the intricate demands of human intent necessitates a more dynamic and adaptive mechanism, giving rise to the critical role of reinforcement learning, particularly Reinforcement Learning from Human Feedback (RLHF), in modern LLM development to address the pervasive 'alignment problem' ?.

The initial motivation for RLHF stemmed directly from this need to imbue LLMs with human values and instruction-following capabilities. Pioneering works by ? and ? introduced the core concept of training a reward model on human preference comparisons, then using this model to fine-tune the LLM policy via reinforcement learning algorithms such as Proximal Policy Optimization (PPO). This approach was famously scaled by ? to train InstructGPT, demonstrating its power in enabling LLMs to follow instructions and exhibit helpful, harmless, and honest behavior. However, PPO-based RLHF proved notoriously complex, unstable, and computationally intensive, requiring careful hyperparameter tuning and the coordination of multiple large models ?. This complexity immediately motivated innovations like Direct Preference Optimization (DPO) by ?, which re-parameterized the RLHF objective into a simpler classification loss, thereby eliminating the need for an explicit reward model and complex reinforcement learning, offering substantial computational and stability advantages.

Despite these foundational advancements, the integration of proxy reward models introduced new, critical challenges, most notably "reward hacking" and "overoptimization" ?. This phenomenon occurs when LLMs exploit flaws or spurious correlations within the learned reward model, leading to misaligned or undesirable outputs despite achieving high reward scores, rather than genuinely improving alignment with human intent. For instance, ? empirically demonstrated that superficial correlations, such as a bias towards longer responses, could account for a significant portion of perceived RLHF improvements, revealing the non-robustness of early reward models. These issues highlighted that the fidelity and robustness of the reward signal itself were paramount, directly motivating a diverse array of research into making reward models more reliable and policies less exploitative. While specific mitigation techniques are detailed in Section 5.2, their emergence was a direct response to these initial alignment failures.

The practical hurdles of scaling RLHF to ever-larger models and diverse applications also presented significant motivations for innovation. The prohibitive cost and time associated with collecting high-quality human feedback spurred the exploration of Reinforcement Learning from AI Feedback (RLAIF), where large language models themselves gen-

erate preference labels, demonstrating comparable performance to human-labeled RLHF ?. However, RLAIF introduced its own set of challenges, such as the potential for AI-generated feedback to exhibit biases, like the "verbosity bias" where LLM judges disproportionately prefer longer responses, even if not qualitatively superior ?. This further motivated the development of more sophisticated AI feedback mechanisms, such as Hybrid RLAIF (HRLAIF) by ?, which aims to balance helpfulness and harmlessness by improving AI's accuracy in discerning correctness. The sheer computational burden and diminishing returns observed in scaling RLHF ? also motivated efforts towards greater efficiency, including methods like Proxy-RLHF ?, which decouples generation and alignment for parameter efficiency, and theoretical work on optimal design for reward modeling to minimize human labeling costs ?.

Beyond technical flaws, the widespread application of RLHF brought to light broader, often undesirable, consequences that impacted the trustworthiness and utility of aligned LLMs. Researchers observed a significant trade-off where RLHF, while improving instruction following, often reduced output diversity, potentially limiting creativity ???. The "alignment tax," or catastrophic forgetting of pre-trained capabilities, became a notable concern, where models lost foundational knowledge in the pursuit of alignment ?. Furthermore, the "objective mismatch" problem, where proxy reward models failed to truly capture nuanced human values, led to issues like excessive refusals or "laziness" in LLM responses ???. The security and privacy implications also emerged as critical motivations: RLHF-trained models were found to be vulnerable to data poisoning attacks ? and "jail-breaking" techniques ?, while memorization of sensitive training data became a concern ?. These unintended consequences underscored the inherent complexity of aligning powerful AI systems, prompting a critical re-evaluation of the sociotechnical limits of AI alignment and safety through RLHF ?.

As LLMs became more sophisticated and their applications diversified, the need for multi-objective alignment and the ability to handle diverse, potentially conflicting, human preferences became paramount. Single-scalar reward models proved insufficient for capturing the full spectrum of human values, such as simultaneously optimizing for help-

fulness, safety, and conciseness ???. This motivated the development of frameworks like MaxMin-RLHF ? to align with diverse user groups, and Pareto-Optimal Preference Learning (POPL) ? to learn policies optimal for distinct hidden context groups. The critical helpfulness-safety trade-off, where overly safe models might become unhelpful, was a major driver for methods like Equilibrate RLHF ?, which uses a fine-grained data-centric approach and adaptive message-wise alignment to achieve a better balance. These advancements reflect a maturation of the alignment goal, moving from simplistic optimization to a more holistic and nuanced understanding of human intent.

In conclusion, the journey of aligning LLMs with human intent through reinforcement learning is a dynamic and multifaceted endeavor. The initial motivations stemmed from the inherent limitations of supervised learning in capturing subjective human preferences and ethical considerations. However, the deployment of foundational RLHF techniques quickly revealed a new set of challenges, including reward model imperfections, scalability issues, and unintended consequences like reduced diversity and the "alignment tax." These persistent problems, coupled with the growing need for multi-objective and robust alignment, continue to drive research into more sophisticated reward modeling, efficient feedback mechanisms, and theoretically grounded approaches. The ongoing quest is to ensure LLMs are not only powerful but also trustworthy, safe, and truly aligned with humanity's best interests, thereby enhancing their utility and trustworthiness in real-world applications ???.

1.3 Scope and Structure of the Review

This literature review provides a comprehensive and critically analyzed overview of Reinforcement Learning (RL) applications in Natural Language Processing (NLP), tracing the field's evolution from foundational concepts to cutting-edge advancements in large language model (LLM) alignment and capabilities. Our scope is primarily focused on text-based NLP tasks, including but not limited to language generation, dialogue systems, summarization, machine translation, and complex instruction following. We delve into the integration of various RL methodologies, particularly Deep Reinforcement Learn-

ing (DRL) and Reinforcement Learning from Human Feedback (RLHF), examining their theoretical underpinnings, practical implementations, and the challenges they address. Crucially, this review deliberately excludes areas such as speech processing, pure computer vision tasks, or robotic control where language interaction is minimal or not the primary focus, ensuring a concentrated analysis of RL's impact on textual language understanding and generation. Furthermore, while acknowledging the breadth of RL, our analysis prioritizes policy-based methods and preference learning paradigms that have proven central to modern LLM alignment, with less emphasis on traditional value-based approaches not directly applied to language generation.

The review is structured to offer a pedagogical progression, beginning with the foundational principles that enabled the framing of language tasks as sequential decision-making problems. Section 2, "Foundational Concepts: Bridging RL and Language," establishes this groundwork by introducing core RL algorithms like policy gradients ?? and the conceptual shift of viewing structured prediction as a search problem, which was vital for optimizing non-differentiable NLP metrics ?. This sets the stage for understanding the subsequent integration of deep learning.

Building on these foundations, Section 3, "Deep Reinforcement Learning for Direct Language Generation," explores early applications of DRL in tasks such as dialogue generation ? and abstractive summarization ?. This section highlights how DRL addressed limitations of supervised learning, such as exposure bias and the inability to optimize for long-term, holistic quality metrics, marking a significant step towards more coherent and human-like generated text.

A pivotal intellectual shift is then examined in Section 4, "The Rise of Reinforcement Learning from Human Feedback (RLHF)." This section details how RLHF revolutionized LLM alignment by leveraging human preferences to train reward models and fine-tune policies, moving beyond proxy metrics to directly optimize for complex, subjective objectives like helpfulness and safety ???. It critically compares key algorithms such as Proximal Policy Optimization (PPO) ? and the more recent Direct Preference Optimization (DPO) ?, which streamlines the alignment process by eliminating the explicit reward model.

Section 5, "Advanced Feedback Mechanisms and Robust Alignment," extends this discussion by exploring sophisticated advancements in feedback generation and utilization. This includes Reinforcement Learning from AI Feedback (RLAIF) ?, and Constitutional AI ?, which aim to scale alignment by leveraging AI-generated preferences. It also addresses critical challenges like reward model imperfections, overoptimization, and the development of interpretable and multi-objective reward modeling to ensure more robust and nuanced alignment ?.

The review then shifts to specialized applications in Section 6, "RL for Enhanced LLM Capabilities and Specialized Applications," showcasing RL's versatility beyond general conversational alignment. This includes enhancing LLM reasoning for mathematical problem-solving ?, enabling tool use and agentic behavior ?, and even automating reward design for external agents, demonstrating RL's role in developing more intelligent and adaptable AI agents.

Finally, Section 7, "Efficiency, Scalability, and Responsible AI in RL for Language Processing," addresses critical practical and ethical considerations. It covers advancements in training efficiency, such as parameter-efficient RLHF ?, dynamic inference-time alignment, and the paramount importance of privacy, security, and robustness in aligned LLMs. This section also provides a critical evaluation of the sociotechnical limits of AI alignment, emphasizing the need for responsible development and robust evaluation to ensure powerful RL-driven LLMs are safe, fair, and trustworthy.

The review concludes in Section 8 by synthesizing these advancements, identifying unresolved theoretical gaps and practical challenges, and outlining promising future directions for research. By following this structured narrative, readers will gain a deep, coherent understanding of RL's transformative impact on NLP, its current state, and the exciting avenues for its continued evolution.

2 Foundational Concepts: Bridging RL and Language

2.1 Core Reinforcement Learning Algorithms for Sequential Decision Making

The inherent sequential nature of language generation, where each token choice influences subsequent possibilities and the overall quality of the generated text, necessitates a decision-making framework capable of optimizing for long-term, cumulative rewards. Reinforcement Learning (RL) provides this framework, training an agent to learn a policy that directly maps states to actions to maximize expected future returns. This section introduces the foundational RL algorithms—policy gradient methods and actor-critic architectures—that form the bedrock for sequential decision-making in language processing, laying the groundwork for more advanced applications.

Policy gradient methods represent a fundamental class of RL algorithms that directly optimize the parameters of a stochastic policy, $\pi_\theta(a|s)$, which defines the probability of taking action a in state s . The core idea is to estimate the gradient of the expected return with respect to the policy parameters θ and update θ in the direction of increasing return. The seminal REINFORCE algorithm, introduced by Williams1992, is a Monte Carlo policy gradient method that estimates this gradient by sampling complete trajectories. For each trajectory, the policy parameters are updated proportionally to the product of the gradient of the log-probability of the taken actions and the total return received from that point onward. Intuitively, the log-derivative trick, $\nabla_\theta \log \pi_\theta(a|s) = \frac{\nabla_\theta \pi_\theta(a|s)}{\pi_\theta(a|s)}$, allows the gradient of the expected return to be expressed as an expectation over the policy’s actions. This means that actions leading to high returns are reinforced (made more probable), while actions leading to low returns are discouraged. Mathematically, the update for a single trajectory $\tau = (s_0, a_0, r_1, s_1, \dots, s_{T-1}, a_{T-1}, r_T)$ is given by $\Delta\theta \propto \sum_t \nabla_\theta \log \pi_\theta(a_t|s_t) G_t$, where G_t is the cumulative discounted return from time step t . This direct optimization approach is particularly advantageous for discrete action spaces, characteristic of token generation in language, and for optimizing non-differentiable, sequence-level rewards, which are common in NLP

evaluation metrics. Sutton2000 further formalized and generalized policy gradient methods with function approximation, laying the theoretical groundwork for their integration with deep neural networks. A significant challenge of pure policy gradient methods like REINFORCE, however, is the high variance of their gradient estimates, which can lead to unstable and inefficient training, especially in complex environments with long horizons. This high variance stems from the reliance on full trajectory returns (G_t), which can fluctuate significantly between samples.

To address the high variance inherent in pure policy gradient methods, actor-critic architectures emerged as a more stable and efficient alternative. These methods combine the strengths of both policy-based (actor) and value-based (critic) RL. The "actor" is responsible for learning the policy, i.e., how to choose actions, while the "critic" learns a value function (e.g., state-value function $V(s)$ or state-action value function $Q(s, a)$) to estimate the expected future reward from a given state or state-action pair. A key innovation of actor-critic methods is the use of bootstrapping, where the critic learns its value estimates by updating them based on other learned value estimates, rather than waiting for full episode returns. This allows for continuous learning and often faster convergence compared to Monte Carlo methods. The critic's value estimates serve as a learned baseline for the actor's policy updates. By subtracting a baseline (typically the state-value function $V(s)$) from the return G_t to form an advantage estimate $A_t = G_t - V(s_t)$, the variance of the policy gradient estimate is significantly reduced without changing its expectation. This is because the baseline term, being independent of the action taken, does not affect the expected gradient but effectively normalizes the reward signal, leading to more stable learning. Early foundational work on actor-critic methods, such as those by Barto1983 and further developed by Konda1999, established their theoretical underpinnings and practical benefits. The evolution of these methods saw the development of algorithms like Asynchronous Advantage Actor-Critic (A3C) and its synchronous variant A2C, which further improved stability and sample efficiency by leveraging parallel environments and more robust advantage estimation techniques. These advancements were crucial in making actor-critic methods practical for complex, high-dimensional problems,

including those in language processing.

In summary, policy gradient methods, exemplified by REINFORCE, and the more stable actor-critic architectures, represent the fundamental algorithmic toolkit for training agents in sequential decision-making tasks critical to language processing. They enable direct optimization for long-term, non-differentiable rewards, addressing key limitations of supervised learning. While pure policy gradients offer conceptual simplicity, actor-critic methods provide a crucial advancement in stability and efficiency through variance reduction and bootstrapping. These foundational algorithms, continuously refined and adapted, remain indispensable for developing sophisticated and human-aligned language AI. The principles established here, particularly the actor-critic paradigm, were later extended and refined into powerful algorithms like Proximal Policy Optimization (PPO), which became central to modern large language model alignment through Reinforcement Learning from Human Feedback (RLHF), a topic that will be explored in detail in Section 4.

2.2 Structured Prediction as a Search Problem

Traditional approaches to structured prediction in Natural Language Processing (NLP), such as parsing or sequence labeling, often relied on complex graphical models or dynamic programming algorithms to find the optimal global structure. However, these methods frequently struggled with the inherent non-local dependencies and the mismatch between local training objectives and global evaluation metrics. This challenge spurred a paradigm shift, reframing structured prediction as a sequential decision-making process, akin to a search problem, where a policy learns to construct the output incrementally. This perspective was crucial for applying reinforcement learning (RL)-like thinking to NLP before the widespread adoption of deep reinforcement learning.

The conceptual groundwork for learning policies that make sequential decisions can be traced back to foundational work in reinforcement learning. Williams1992 introduced the REINFORCE algorithm, a seminal policy gradient method, demonstrating how to directly optimize a policy’s parameters to maximize expected rewards without requiring a

differentiable model of the environment. Building on this, Sutton2000 further elaborated on policy gradient methods with function approximation, providing a robust theoretical framework for learning policies in complex, high-dimensional state spaces. While these works were general to RL, they established the core algorithmic toolkit for learning decision-making policies, which would later influence structured prediction.

In the context of NLP, early explorations began to frame structured output generation as a sequence of choices. For instance, Ranzato2007 explored learning policies for sequence labeling in computer vision, demonstrating how to construct structured outputs step-by-step by learning a policy that makes local decisions. This work highlighted the potential of moving beyond static, one-shot predictions to a dynamic, generative process.

A pivotal contribution that formalized this conceptual bridge for NLP was the "learning to search" (L2S) framework, prominently articulated by Daume2009. This framework reframes complex structured prediction tasks, such as dependency parsing or machine translation, as a series of local classification decisions made by a learned policy. Instead of relying on a pre-defined search algorithm, L2S trains a policy to guide a search process through the space of possible output structures, aiming to find the globally optimal one. This is often achieved through imitation learning, where the policy is trained to mimic an expert's decisions during a search, effectively reducing structured prediction to a sequence of supervised classification problems. The core idea is to learn a policy that makes decisions at each step of the output construction, thereby addressing the sequential nature of linguistic outputs and the need for global optimization beyond local token-level accuracy.

The significance of the L2S framework and its broader implications for structured learning in NLP were further consolidated by comprehensive surveys. Chang2015 provided a detailed overview of structured learning techniques for NLP, including the L2S paradigm, solidifying its importance as a method for tackling complex linguistic tasks. This work underscored how L2S offered a flexible and powerful way to handle the interdependencies inherent in structured outputs, moving beyond the limitations of independent local predictions.

Despite its innovative approach to global optimization, the L2S framework, particu-

larly when relying on imitation learning, faced inherent limitations. The primary challenge was "exposure bias," where the learned policy, trained on expert demonstrations, might encounter states during inference that were never seen during training, leading to compounding errors. Furthermore, while L2S aimed for global optimization, it often did so indirectly by mimicking an expert, rather than directly optimizing for non-differentiable, task-specific evaluation metrics. These unresolved tensions, particularly the exposure bias and the indirect nature of optimization, laid the groundwork for the subsequent advent of deep reinforcement learning methods that sought to directly optimize for long-term, non-differentiable rewards in sequence generation tasks.

2.3 Optimizing Non-Differentiable Metrics in NLP

Traditional approaches to training sequence generation models in Natural Language Processing (NLP) predominantly rely on Maximum Likelihood Estimation (MLE). MLE optimizes the probability of generating the correct next token given previous tokens and the ground truth, proving effective for learning language distributions. However, this token-level optimization suffers from inherent limitations when the ultimate goal is to produce outputs that score well on holistic, non-differentiable, task-specific evaluation metrics such as BLEU for machine translation or ROUGE for summarization, or more broadly, human-aligned quality judgments ?. This fundamental discrepancy between the training objective (local token accuracy) and the evaluation objective (global sequence quality) often leads to issues like "exposure bias" ?. Exposure bias occurs because models are trained on ground-truth prefixes but are forced to generate subsequent tokens based on their own potentially erroneous outputs during inference, leading to a compounding of errors and a divergence from the desired output distribution.

Early research highlighted the critical need to bridge this gap. While recurrent neural networks (RNNs) demonstrated promise for sequence transduction tasks under supervised learning ?, their performance was fundamentally constrained by the MLE objective. The desire for direct metric optimization became evident with methods like Minimum Error Rate Training (MERT) ?. MERT discriminatively tuned parameters for statistical

machine translation systems to maximize BLEU scores. Critically, MERT operated by performing a line search in parameter space to directly optimize the non-differentiable metric, demonstrating the feasibility and importance of such direct optimization. However, MERT was computationally intensive, specific to feature-based translation models, and not readily adaptable to the end-to-end gradient-based training of deep neural networks. Concurrently, frameworks for structured prediction, such as "learning to search," conceptualized complex NLP tasks like parsing or sequence labeling as a sequence of local classification decisions ?. While these methods often leveraged imitation learning to guide a search process, they underscored the sequential decision-making paradigm inherent in language generation and the need for global optimization beyond local accuracy. These precursors collectively identified the core problem: MLE's inability to directly optimize for human-aligned, non-differentiable metrics and the practical limitations of pre-deep learning discriminative methods.

This persistent challenge motivated the adoption of reinforcement learning (RL), which offers a powerful paradigm to directly optimize non-differentiable objectives without requiring explicit gradient information from the reward function itself. This capability enables models to generate text that is more aligned with human judgments of quality. A pivotal conceptual breakthrough was the application of policy gradient methods, particularly REINFORCE ?, to sequence-to-sequence learning. This approach reframed text generation as a sequential decision-making process, where each generated token is an action, and the entire generated sequence receives a scalar reward based on a non-differentiable metric (e.g., BLEU, ROUGE). By sampling sequences from the model's current policy and using the resulting metric score as a reward signal, REINFORCE allowed models to be trained to directly maximize these sequence-level metrics. This inherently mitigated exposure bias, as the model learned from its own generated sequences rather than solely from ground truth. Subsequent work also explored actor-critic algorithms for sequence prediction, aiming to reduce the high variance often associated with pure policy gradient methods and improve training stability ?.

However, the practical application of these early RL methods for sequence generation

was not without its challenges. While conceptually sound, policy gradient methods often suffered from high variance and sample inefficiency, requiring extensive exploration of the vast action space (vocabulary) and long sequences. The computational burden of sampling numerous sequences for reward estimation and gradient calculation, especially for long sequences and large vocabularies, posed a significant practical hurdle, impacting training time and memory consumption ?. This inherent inefficiency limited the scalability and widespread adoption of RL for direct metric optimization in complex NLP tasks, even as the conceptual framework proved its merit. Despite these practical difficulties, the introduction of RL provided the foundational mechanism to bridge the gap between token-level training and sequence-level evaluation, laying the groundwork for later deep RL applications in language generation.

3 Deep Reinforcement Learning for Direct Language Generation

3.1 DRL for Dialogue Generation

The initial wave of dialogue system development largely relied on supervised learning, which, despite its efficacy in predicting next tokens or turn-level responses, inherently struggled to optimize for the holistic, long-term qualities crucial for engaging conversations, such as coherence, naturalness, and overall user satisfaction ?. This limitation stemmed from issues like exposure bias, where models trained on ground-truth sequences diverge from optimal paths during self-generated inference, leading to generic, repetitive, or short-sighted responses. Deep Reinforcement Learning (DRL) emerged as a transformative paradigm, framing dialogue generation as a sequential decision-making process. This allowed agents to learn optimal policies by interacting with an environment (e.g., a simulated user) and optimizing for cumulative rewards over multiple turns, thereby directly addressing the shortcomings of purely supervised approaches. The conceptual foundation for optimizing non-differentiable, sequence-level metrics, crucial for DRL in language gen-

eration, was established in earlier works on structured prediction, as discussed in Section 2.

A seminal contribution to the application of DRL in dialogue generation was the work by ?. This research pioneered the use of deep reinforcement learning to train neural dialogue agents, moving beyond simple next-utterance prediction to optimize for long-term conversational objectives. Their approach explicitly aimed to improve metrics such as ease of answering, information flow, and overall coherence, which are notoriously difficult to capture with conventional turn-level supervised losses. By framing dialogue as a Markov Decision Process, the authors employed policy gradient methods, specifically REINFORCE, to train an agent that could learn from interactions with a simulated user. The agent received sequence-level rewards designed to encourage desirable conversational traits, such as generating diverse and informative responses while avoiding repetition. This marked a significant departure from previous methods, demonstrating DRL’s potential to mitigate exposure bias and foster more engaging, goal-oriented dialogues. However, this pioneering work, like many early applications of pure policy gradient methods, faced inherent challenges, including high variance during training and the significant difficulty of designing effective, non-sparse reward functions for complex, multi-turn interactions.

Following ?, subsequent research further explored the application of DRL to dialogue, particularly focusing on refining reward mechanisms and improving training stability. Early efforts often distinguished between open-domain and task-oriented dialogue systems. In task-oriented settings, DRL was applied to learn optimal dialogue policies for achieving specific goals, such as booking flights or making reservations. Here, rewards could be more explicitly defined based on task success, user satisfaction (often simulated), and efficiency metrics like turn count ?. For instance, ? utilized DRL to learn dialogue policies for goal-oriented conversational agents, demonstrating improvements in task completion rates and user satisfaction by optimizing a policy network. These systems often leveraged simulated user models to provide a rich, interactive training environment, allowing agents to explore diverse dialogue strategies and learn from long-term consequences. In open-domain dialogue, where the goal is more about engaging and coherent conversation rather

than task completion, reward design became more abstract and challenging. Researchers experimented with various forms of intrinsic and extrinsic rewards, including those based on diversity, informativeness, and avoiding repetition. Some approaches explored adversarial training, where a discriminator network provided a "naturalness" reward signal to the generator, pushing it to produce more human-like responses ⁷. These methods aimed to overcome the limitations of simple n-gram based metrics and the difficulty of obtaining explicit human feedback for every interaction.

Despite these pioneering efforts, several fundamental challenges persisted in the early applications of DRL for dialogue generation. Reward sparsity remained a significant hurdle, as meaningful positive feedback often materialized only at the end of a multi-turn conversation, making it difficult for agents to attribute credit to individual actions. The vast, discrete action space of language also exacerbated the exploration-exploitation dilemma, making it computationally expensive for agents to discover optimal dialogue policies efficiently. Furthermore, the high variance associated with pure policy gradient methods often led to unstable training. These challenges underscored the critical need for more robust reward engineering and efficient learning from feedback. While initial DRL for dialogue often relied on hand-crafted or simple simulated rewards, the field progressively recognized the importance of more nuanced feedback mechanisms. For instance, the difficulty in defining explicit, non-sparse rewards for long, multi-turn interactions directly motivated later advancements in learning from conversation-level preferences. ⁸, for example, formalizes multi-turn preference optimization, where feedback is provided on *entire conversational trajectories* rather than individual turns. This approach, while a more recent development in the RLHF paradigm (as discussed in Section 4), directly addresses the long-standing problem of reward sparsity and the need to optimize for long-term conversational quality that was first highlighted in pioneering DRL for dialogue. Similarly, as reward models became more sophisticated, the challenge of "reward model overoptimization" emerged, where maximizing a proxy reward could lead to a decrease in actual human-judged quality. This problem, exemplified by work like ⁹ in dialogue generation, highlights the continuous complexity of aligning learned rewards with true human

preferences, a challenge that traces its roots back to the initial difficulties of designing effective reward functions for DRL-driven dialogue agents.

In summary, the pioneering applications of deep reinforcement learning to dialogue generation marked a crucial shift from turn-level accuracy to optimizing for long-term conversational quality and user satisfaction. By framing dialogue as a sequential decision-making problem, DRL enabled agents to learn adaptive policies from interactive feedback, moving beyond the generic responses of purely supervised models. While initial methods faced significant hurdles related to reward design, sparsity, and training stability, these early explorations laid the groundwork for future advancements in more robust feedback mechanisms and sophisticated policy optimization, continually pushing towards more engaging, coherent, and goal-oriented interactive language systems.

3.2 DRL for Abstractive Summarization and Machine Translation

Traditional sequence-to-sequence models for abstractive summarization and machine translation, primarily trained with Maximum Likelihood Estimation (MLE), often suffer from "exposure bias" and a fundamental mismatch between their training objective (token-level accuracy) and the holistic, non-differentiable metrics used for evaluation (e.g., ROUGE for summarization, BLEU for machine translation). Deep Reinforcement Learning (DRL) emerged as a powerful paradigm to address these limitations by directly optimizing for these sequence-level quality metrics, leading to more fluent, coherent, and human-like generated text.

Early work laid the groundwork for applying reinforcement learning to sequence generation. ? pioneered the use of policy gradient methods for sequence-to-sequence learning, demonstrating how models could be trained to directly optimize non-differentiable metrics like BLEU. This approach allowed the model to explore the output space and learn to generate sequences that were globally better, rather than just locally accurate. The pervasive issue of exposure bias, where models trained on ground truth sequences struggle when exposed to their own generated (and potentially erroneous) tokens during inference, was

a key motivator for this shift. ?, in the context of dialogue generation, further highlighted this problem and showcased how policy gradient methods could mitigate it by allowing the model to learn from its own generated sequences, thereby improving long-term coherence and relevance.

For abstractive summarization, DRL has been instrumental in enhancing the quality of generated summaries beyond mere extractive approaches. ? introduced a deep reinforced model for abstractive summarization, leveraging an actor-critic architecture to directly optimize ROUGE scores. This method combined supervised pre-training with DRL fine-tuning, allowing the model to learn to generate more coherent and informative summaries by rewarding it for higher ROUGE scores, which are indicative of content overlap and fluency. Building on this, ? (Deep Reinforcement Learning for Abstractive Summarization) further explored DRL techniques to improve abstractive summarization, often focusing on refining reward functions and training stability to produce more factually consistent and readable outputs.

In the domain of neural machine translation (NMT), DRL has similarly been employed to bridge the gap between training and evaluation. ? proposed an actor-critic algorithm specifically for sequence prediction, which is highly applicable to NMT. This method aimed to reduce the high variance associated with pure policy gradient methods by introducing a critic network to estimate the value function, thereby providing more stable and efficient learning signals. This stability is crucial for complex tasks like machine translation, where the output space is vast. Furthermore, ? (Reinforcement Learning for Neural Machine Translation with a Diversified Reward) explicitly applied reinforcement learning to NMT, designing diversified reward functions that not only considered BLEU scores but also encouraged diversity in translations, addressing the issue of generic or repetitive outputs. Another facet of DRL’s utility in NMT was demonstrated by ? (A Reinforcement Learning Approach to Improve the Robustness of Neural Machine Translation), where policy gradient methods were used to train an agent to generate adversarial examples, which in turn improved the robustness of the NMT model. This application showcases DRL’s capacity to enhance the overall reliability and quality of generated con-

tent, a critical step for real-world utility.

In conclusion, the application of DRL, particularly through policy gradients and actor-critic models, marked a significant advancement in abstractive summarization and machine translation. By directly optimizing for holistic, non-differentiable metrics like ROUGE and BLEU, DRL successfully addressed the limitations of MLE, such as exposure bias, leading to the generation of more fluent, coherent, and human-like text. Despite these successes, challenges persist, including the inherent high variance of RL training, the sample inefficiency for complex language tasks, and the difficulty of designing effective and generalizable reward functions that truly capture nuanced aspects of text quality beyond n-gram overlap. These unresolved tensions continue to drive research towards more stable, efficient, and sophisticated DRL paradigms for natural language generation.

3.3 Actor-Critic Methods and Exposure Bias Mitigation

The application of Deep Reinforcement Learning (DRL) to sequence generation tasks, particularly in natural language processing, has been driven by the need to overcome fundamental limitations of traditional supervised learning, most notably the "exposure bias" problem and the challenge of optimizing for non-differentiable, sequence-level metrics. Exposure bias arises when models are trained on ground truth sequences using maximum likelihood estimation (MLE) but are then exposed to their own generated tokens during inference, leading to a compounding of errors and degraded performance ?. DRL offers a paradigm shift by framing sequence generation as a sequential decision-making process, where an agent learns a policy to generate tokens that maximize a long-term, task-specific reward.

Early efforts to apply reinforcement learning to structured prediction and sequence generation, such as the work by ?, demonstrated the potential of policy gradient methods like REINFORCE to optimize for sequence-level objectives. Similarly, ? pioneered the use of policy gradient DRL for dialogue generation, explicitly highlighting the pervasive exposure bias problem in this context. ? further applied policy gradient methods to sequence-to-sequence learning, aiming to directly optimize non-differentiable metrics and

mitigate exposure bias. While these pure policy gradient approaches offered a direct way to optimize for desired outcomes, they often suffered from high variance during training, making learning unstable and sample-inefficient, as detailed in foundational works on policy gradient methods with function approximation [1]. The conceptual framework of "learning to search" [2] also provided a precursor, framing structured prediction as a sequence of local decisions, which DRL later leveraged for more complex generative tasks.

To address the high variance inherent in pure policy gradient methods and enhance training stability, actor-critic algorithms emerged as a crucial advancement. These methods combine a policy network (the "actor") that selects actions (tokens) with a value network (the "critic") that estimates the expected return of states or state-action pairs, thereby providing a lower-variance baseline for the policy gradient. [3] introduced an actor-critic algorithm specifically designed for sequence prediction, demonstrating its effectiveness in reducing variance and improving the stability of DRL training for tasks like abstractive summarization. This approach allowed for more efficient learning by leveraging the critic's estimate to guide the actor's updates. Following this, [4] successfully applied an actor-critic model to abstractive summarization, optimizing directly for ROUGE scores and further showcasing the method's ability to mitigate exposure bias by training on the actual evaluation metric rather than token-level accuracy. Other works, such as [5], also explored DRL with diversified rewards for neural machine translation, contributing to the broader effort of using RL to optimize for specific generation qualities beyond MLE.

The evolution of DRL for language generation also saw the development of various strategies to explicitly mitigate exposure bias. Beyond the inherent benefit of optimizing sequence-level rewards, methods often combined supervised pre-training with DRL fine-tuning, allowing the model to first learn basic generation capabilities from ground truth before refining its policy with RL to handle its own generated outputs. This hybrid approach, often employing actor-critic variants, became a standard practice. Comprehensive surveys such as [6, 7, 8] consistently highlight exposure bias as a central challenge in DRL for natural language generation and extensively review the role of actor-critic methods and other DRL strategies in addressing it.

More recently, the principles of actor-critic methods have been scaled and refined in the context of large language models (LLMs) through Reinforcement Learning from Human Feedback (RLHF). Pioneering works like ? and ? introduced the concept of training a reward model from human preferences, which then serves as the critic to fine-tune the LLM (the actor) using policy optimization algorithms like Proximal Policy Optimization (PPO), an advanced actor-critic variant. This paradigm, famously scaled in ? to create InstructGPT, represents a sophisticated approach to aligning LLMs with complex human instructions and values. By optimizing for human-preferred outcomes, RLHF inherently addresses a form of exposure bias, where the model learns to generate outputs that are not just grammatically correct but also helpful, harmless, and honest, even when exposed to its own generated content during the iterative refinement process ??. This demonstrates a significant leap in using actor-critic principles to mitigate the mismatch between training objectives and desired generation quality. Furthermore, preference-based RL, as seen in ?, leverages LLMs to reason from human text prompts, further refining reward learning and policy shaping.

In conclusion, the journey from pure policy gradient methods to sophisticated actor-critic algorithms and RLHF has been pivotal in making DRL practical for complex language generation tasks. Actor-critic methods significantly reduced variance and improved training stability, directly addressing a key limitation of earlier DRL applications. This, coupled with direct optimization of sequence-level rewards, has been instrumental in mitigating exposure bias. However, challenges persist, including the complexity of designing effective and generalizable reward functions, the sample inefficiency inherent in RL, and ensuring true alignment with nuanced human preferences in large-scale RLHF applications. Future research continues to explore more robust and efficient ways to leverage these methods for high-quality and reliable language generation.

4 The Rise of Reinforcement Learning from Human Feedback (RLHF)

4.1 Core RLHF Paradigm: Reward Modeling and Policy Optimization

The evolution of large language models (LLMs) from sophisticated text generators to capable instruction followers and value-aligned agents marks a pivotal shift, largely driven by the Reinforcement Learning from Human Feedback (RLHF) paradigm. Prior deep reinforcement learning (DRL) methods for sequence generation, while innovative in optimizing non-differentiable metrics, often grappled with the inherent challenge of designing effective, hand-crafted reward functions that truly captured the nuanced, subjective aspects of human quality and preferences ??? . These proxy rewards frequently failed to align model outputs with complex human judgments, leading to outputs that were grammatically correct but semantically or pragmatically misaligned.

The RLHF paradigm fundamentally addressed these limitations by introducing a multi-stage process that learns directly from human preferences. This approach, initially explored in foundational works such as ? and ?, begins with the collection of human preference data, typically through pairwise comparisons of different model outputs for a given prompt. For instance, annotators might be asked to select which of two generated summaries is better, or which dialogue response is more helpful. This human preference data is then used to train a separate reward model (RM), often a neural network itself, to predict human preferences. The reward model learns to assign a scalar score to any given text sequence, effectively encoding the complex and subjective aspects of human judgment that were previously intractable to define programmatically ?? . This step is crucial as it transforms the qualitative, comparative human feedback into a quantitative, differentiable reward signal, thereby overcoming the bottleneck of hand-crafted or proxy rewards.

Following the training of the reward model, the second stage involves policy optimiza-

tion. Here, the pre-trained language model, which serves as the policy, is fine-tuned using reinforcement learning algorithms, most commonly Proximal Policy Optimization (PPO) [1]. The reward model, now acting as the objective function, provides a scalar reward for the language model’s generated outputs. The language model is then updated to maximize this learned reward, effectively learning to generate text that is highly rated by the reward model, and by extension, aligned with human preferences [2]. This iterative process allows the LLM to refine its behavior, moving beyond mere statistical text generation to actively produce responses that are helpful, harmless, and honest, as demonstrated by the significant advancements in models like InstructGPT and its successor ChatGPT [3]. The success of this approach is evident in its ability to enable LLMs to follow complex instructions and align with human values, a capability largely absent in earlier models that relied solely on pre-training and supervised fine-tuning. For example, GPT-3.5’s ability to provide appropriate responses to diverse user queries was largely attributed to this RLHF finetuning process [4].

This paradigm shift has not only enhanced the instruction-following capabilities of LLMs but has also opened avenues for addressing more complex alignment challenges. However, the RLHF paradigm is not without its limitations. The quality and scalability of human preference data remain critical concerns, as biased or insufficient data can lead to a reward model that misrepresents true human preferences, potentially resulting in "reward hacking" where the model optimizes for the reward signal without truly achieving the desired behavior. The persistent issue of hallucinations, where LLMs generate factually incorrect or nonsensical information, continues to be a challenge even with RLHF, as highlighted by recent analyses of models like ChatGPT and GPT-4 [5]. Recent efforts, such as those in autonomous driving, have begun to integrate RLHF with multimodal LLMs, leveraging iterative refinement loops and symmetric performance metrics (e.g., BLEU, ROUGE) to explicitly target hallucination reduction and optimize computational efficiency [6]. This suggests an ongoing evolution where RLHF is combined with other techniques to enhance robustness and address specific failure modes. Despite these advancements, ensuring robust, scalable, and truly aligned behavior across diverse

and complex contexts remains a key area of ongoing research, particularly as LLMs are deployed in safety-critical applications.

4.2 Proximal Policy Optimization (PPO) for LLM Alignment

Proximal Policy Optimization (PPO) stands as the predominant algorithm for the policy optimization stage in Reinforcement Learning from Human Feedback (RLHF), serving as a cornerstone for aligning Large Language Models (LLMs) with human preferences ???. PPO's widespread adoption stems from its balance of stability and sample efficiency, achieved through a clipped surrogate objective and multiple epochs of optimization on collected data, which helps prevent excessively large policy updates. However, applying PPO to the vast and complex action spaces of LLMs presents unique challenges, including training instability, hyperparameter sensitivity, and the intricate coordination of multiple models ?.

To address these inherent difficulties, significant advancements have been made in refining PPO for LLM alignment. zheng2023c98 meticulously dissected PPO's inner workings and proposed **PPO-max**, a carefully calibrated collection of effective PPO implementations designed to ensure stable and efficient policy model training. Their work identified policy constraints as a key factor for effective implementation and introduced action space modeling metrics (e.g., perplexity, response length, KL divergence between policy and SFT model) as more informative indicators of stability than traditional reward and loss functions. This foundational work provided a robust PPO variant, making the RLHF pipeline more accessible and reliable.

Despite these algorithmic refinements, PPO-based RLHF remains susceptible to "reward hacking" or "overoptimization," where the policy exploits imperfections in the learned reward model rather than genuinely improving alignment. singhal2023egk empirically demonstrated that spurious correlations, such as length bias, can significantly drive reported RLHF improvements, revealing that reward models often learn shallow features. This issue is further explored by rafailov2024ohd, who established scaling laws for reward model overoptimization in Direct Alignment Algorithms (DAAs), showing that

performance degradation can occur very early in training, even in methods that bypass explicit reward models, underscoring the fragility of preference-based optimization.

A wave of research has focused on mitigating these reward model imperfections and their impact on PPO training. [yu20249l0](#) enhanced reward model quality and interpretability by leveraging self-generated critiques, thereby providing a more robust signal for PPO to optimize. For scenarios involving composite reward models, [moskovitz2023slz](#) introduced constrained RLHF with dynamic proxy points, preventing PPO from overoptimizing individual reward components beyond their effective range. Directly modifying the reward signal for PPO, [fu2025hl3](#) proposed **Preference As Reward (PAR)**, a novel reward shaping technique that applies a sigmoid function to centered rewards. This method is theoretically grounded to stabilize critic training and minimize policy gradient variance, effectively mitigating reward hacking. Complementing this, [miao2025ox0](#) introduced **Energy loss-aware PPO (EPPO)**, a mechanistic approach that penalizes the increase in "energy loss" within the LLM's final layer, addressing reward hacking from an internal model dynamics perspective.

Further efforts to make PPO robust against reward model uncertainties include [zhai20238xc](#)'s **Uncertainty-Penalized RLHF (UP-RLHF)**, which uses diverse reward LoRA ensembles to quantify reward uncertainty and penalize rewards based on this estimate, preventing the policy from exploiting out-of-distribution (OOD) regions where the reward model is unreliable. Similarly, [dai2025ygq](#) proposed **Behavior-Supported Policy Optimization (BSPO)**, a value regularization technique that restricts PPO policy iteration to the in-distribution region of the reward model, thereby preventing extrapolation errors. More broadly, [zhang2024esn](#) developed **Adversarial Policy Optimization (AdvPO)**, a distributionally robust optimization for PPO that leverages lightweight uncertainty estimation from last layer embeddings to adversarially search for the most pessimistic reward within a confidence region, making PPO more resilient to reward model inaccuracies without the heavy computational cost of ensembles.

Beyond single-objective overoptimization, PPO for LLMs faces significant challenges in multi-objective alignment and practical deployment. The insufficiency of single re-

ward models for diverse human preferences is highlighted by chakraborty20247ew, who proposed MaxMin-RLHF to ensure fairness and represent minority opinions, an objective that PPO could then optimize. The critical helpfulness-safety trade-off is a major concern, addressed by tan2025lk0 with "Equilibrate RLHF," which includes an Adaptive Message-wise Alignment (AMA) approach that can be integrated with Adaptive PPO (APPO) to selectively emphasize safety-critical segments. xu20242yo introduced **Constrained Generative Policy Optimization (CGPO)** with a "Mixture of Judges" for multi-task learning, developing new primal-type constrained RLHF optimizers (CRPG, CODPO, CRRAFT) that offer robust alternatives or enhancements to PPO in complex multi-objective settings. Furthermore, boldi2024d0s proposed Pareto-Optimal Preference Learning (POPL) to learn a set of reward functions or policies that are Pareto-optimal for distinct hidden context groups, which could subsequently guide PPO optimization. To address sparse and inconsistent human feedback, lai2024ifx introduced ALaRM, a hierarchical reward modeling framework that provides more consistent and fine-grained signals for PPO-like algorithms.

Practical challenges like catastrophic forgetting and training instability in production environments have also driven PPO advancements. hou2024tvy detailed the ChatGLM-RLHF pipeline, offering practical solutions for large-scale LLMs, including strategies for PPO stability, reward debiasing (e.g., bucket-based length balancing), and mitigating catastrophic forgetting by incorporating next-token-prediction loss from SFT data into the RLHF objective. To further combat the "alignment tax" (degradation of pre-trained capabilities), lu202435m proposed **Online Merging Optimizers** (e.g., OnDARE, On-TIES) that integrate model merging into *each* optimization step of RLHF, dynamically steering PPO gradients to boost rewards while preserving foundational capabilities.

In conclusion, PPO remains a central algorithm for LLM alignment, with continuous research dedicated to enhancing its stability, robustness, and applicability to complex, multi-objective scenarios. The evolution from generic PPO to specialized variants like PPO-max, coupled with sophisticated regularization techniques and multi-objective frameworks, underscores a relentless effort to mitigate reward model imperfections, pre-

vent catastrophic forgetting, and navigate the vast action spaces of LLMs, ensuring stable and effective alignment. Future directions will likely continue to explore more adaptive and data-efficient PPO variants, integrating deeper theoretical understandings of model internals and human preferences to achieve truly robust and trustworthy AI systems.

4.3 Direct Preference Optimization (DPO) and Reward-Model-Free Alignment

The traditional Reinforcement Learning from Human Feedback (RLHF) pipeline, while effective for aligning large language models (LLMs) with human preferences, is notoriously complex, computationally intensive, and often unstable, typically involving a separate reward model and a reinforcement learning stage with algorithms like Proximal Policy Optimization (PPO) ?. This complexity spurred the development of more streamlined, reward-model-free alignment approaches, with Direct Preference Optimization (DPO) emerging as a pivotal simplification that democratizes access to preference-based fine-tuning ??.

Direct Preference Optimization (DPO), introduced by rafailov20239ck, fundamentally re-conceptualizes the RLHF objective. Instead of explicitly training a reward model and then using reinforcement learning to optimize the policy, DPO leverages a theoretical insight: the optimal policy for the standard KL-constrained RLHF objective can be derived in closed form by re-parameterizing the reward model. This allows for direct optimization of the language model policy from human preferences using a simple binary cross-entropy loss, effectively eliminating the need for an explicit reward model and complex reinforcement learning algorithms like PPO. DPO’s benefits are substantial, including improved computational efficiency, enhanced training stability, and significantly simplified implementation, making preference-based alignment more accessible.

Beyond DPO, other reward-model-free approaches also aim to simplify the alignment process. hejna2023vyy proposed Contrastive Preference Learning (CPL), which directly learns an optimal policy from regret-based human preferences without an intermediate reward function or traditional RL. CPL leverages the bijection between the optimal ad-

vantage function and the optimal policy’s log-probability in Maximum Entropy RL, transforming the problem into a supervised contrastive objective applicable to general sequential decision-making tasks, including high-dimensional robotics, thereby offering broad applicability beyond just LLMs. Building upon the DPO paradigm, ji2024d5f introduced Self-Play with Adversarial Critic (SPAC), a DPO-like single-timescale algorithm for offline LLM alignment that provides provable convergence guarantees under weak data assumptions, addressing the lack of theoretical robustness in many empirical methods while maintaining computational scalability.

Despite their simplicity, direct alignment algorithms (DAAs) like DPO are not immune to challenges such as reward overoptimization. rafailov2024ohd empirically demonstrated that DAAs exhibit similar performance degradation patterns to classical RLHF, where maximizing the implicit reward beyond a certain point leads to a decline in true human-judged quality. This highlights the need for robust mechanisms even in simplified pipelines. To mitigate such issues, zhai20238xc proposed Uncertainty-Penalized RLHF (UP-RLHF), which augments the objective with an uncertainty regularization term derived from a diverse reward LoRA ensemble, penalizing rewards based on the reward model’s estimated uncertainty. Similarly, miao2025ox0 introduced Energy loss-aware PPO (EPPO), which tackles reward hacking by penalizing the increase in "energy loss" in the LLM’s final layer, offering a novel internal-mechanistic perspective on mitigating policy degeneration. dai2025ygq addressed reward over-optimization via Behavior-Supported Regularization, using value regularization to restrict policy iteration to the in-distribution region of the reward model, thereby preventing extrapolation errors. Furthermore, zhang2024esn proposed Adversarial Policy Optimization (AdvPO) with lightweight uncertainty estimation, a distributionally robust optimization procedure that efficiently quantifies reward uncertainty using only the last layer embeddings of a single reward model, offering a computationally efficient alternative to ensemble methods.

The DPO paradigm has also been extended to address more complex alignment objectives, particularly multi-objective alignment. zhang2024b6u introduced Bi-Factorial Preference Optimization (BFPO), a supervised learning framework that re-parameterizes

a joint RLHF objective for both safety and helpfulness into a single supervised learning loss. This innovation efficiently balances conflicting objectives without the extensive computational and human labor typically required for multi-objective RLHF. tan2025lk0 developed Equilibrate RLHF, which includes an Adaptive Message-wise Alignment (AMA) approach that uses gradient masking to selectively highlight safety-critical segments, adapting DPO (as Adaptive DPO, ADPO) for a more nuanced helpfulness-safety trade-off. xu20242yo proposed Constrained Generative Policy Optimization (CGPO) with a "Mixture of Judges," which includes DPO-like optimizers (Constrained Online DPO, CODPO) to address reward hacking and extreme multi-objective optimization by integrating rule-based and LLM-based judges to identify constraint violations. boldi2024d0s presented Pareto-Optimal Learning from Preferences with Hidden Context (POPL), a reward-model-free approach that leverages lexicase selection to learn a set of Pareto-optimal policies for diverse human preferences without requiring explicit group labels, thereby ensuring fairness and robust personalization.

While DPO and other reward-model-free methods significantly simplify the alignment pipeline, practical challenges persist. The "alignment tax," where fine-tuning for preferences can degrade foundational capabilities, remains a concern. lu202435m introduced Online Merging Optimizers to mitigate this by integrating model merging into each optimization step of RLHF, continuously steering gradients towards maximizing rewards while preserving pre-trained capabilities. This highlights that even with simplified objectives, ensuring comprehensive model performance requires sophisticated regularization and optimization strategies. The ongoing research in this area continues to focus on making preference-based fine-tuning more robust, scalable, and capable of handling complex, multi-faceted alignment objectives, moving towards more reliable and human-aligned LLMs.

4.4 Theoretical Foundations of KL-Constrained Preference Learning

The alignment of large language models (LLMs) with human preferences, predominantly through Reinforcement Learning from Human Feedback (RLHF), necessitates robust theoretical foundations to ensure stability, performance, and a deeper understanding of the underlying optimization problem. A critical aspect of this theoretical grounding involves the use of Kullback-Leibler (KL) divergence constraints, which are instrumental in preventing catastrophic forgetting, maintaining output diversity, and bridging the gap between empirical success and theoretical robustness in LLM alignment.

A pivotal contribution to this field is the formalization of RLHF as a reverse-KL regularized contextual bandit problem by ?. This work provides the first rigorous theoretical analysis of this formulation, offering finite-sample guarantees for offline, online, and hybrid learning settings. By explicitly incorporating a KL-divergence constraint to keep the fine-tuned policy close to a reference model, ? ensures output diversity and fidelity, mitigating issues like reward hacking and the "alignment tax" prevalent in earlier empirical methods. Building upon this, ? further strengthens the theoretical understanding of online KL-regularized reinforcement learning by achieving groundbreaking logarithmic regret bounds for both contextual bandits and Markov Decision Processes. This work provides the first theoretical justification for the superior sample efficiency observed in practical KL-regularized RLHF applications, demonstrating that such regularization fundamentally improves learning efficiency without requiring strong coverage assumptions.

While KL-regularization is crucial for stability, it can also introduce subtle algorithmic biases. ? identifies "preference collapse" as an inherent algorithmic bias in standard KL-regularized RLHF, where the model disproportionately favors dominant preferences, even with an oracle reward model. To address this, they propose Preference Matching (PM) RLHF, which introduces a novel regularizer derived from an ordinary differential equation to provably align LLMs with the preference distribution of the reward model, thereby ensuring that the policy accurately reflects diverse human preferences and maintains output diversity.

The practical implementation of KL-constrained learning often involves iterative algorithms and efficient reward signal generation. ? introduces Reinforced Self-Training (ReST), an iterative, decoupled approach that continuously generates new, higher-quality data from an improving policy and filters it based on a learned reward model. While ReST’s core objective is not explicitly KL-constrained, the reward models it leverages are typically trained using preference data that implicitly or explicitly derive from KL-regularized objectives, highlighting the importance of efficient iterative learning. Extending this, ? proposes DICE, an iterative self-alignment method that bootstraps new preference data using the implicit reward model derived from a DPO-tuned LLM (which itself is based on a KL-regularized objective). This approach enables cost-effective iterative refinement, mitigating catastrophic forgetting and length bias without external feedback. Similarly, ? introduces Self-Exploring Language Models (SELM), an online direct alignment algorithm that uses an optimistically biased, RM-free objective with a KL-divergence constraint to actively explore the response space, preventing models from getting stuck in local optima and promoting guided exploration.

The interplay between online and offline learning, and the robustness of KL-constrained methods, is also a significant area of theoretical inquiry. ? presents Value-Incentivized Preference Optimization (VPO), a unified framework for online and offline RLHF that implicitly handles uncertainty (optimism for online, pessimism for offline) through value function regularization. VPO provides theoretical regret guarantees for both settings, matching standard RL rates while maintaining a DPO-like simplicity, thereby reinforcing the theoretical robustness of KL-constrained approaches in diverse data collection paradigms. Complementing this, ? empirically dissects the performance gap between online and offline alignment algorithms, emphasizing the pivotal role of on-policy sampling in achieving high generative quality, a challenge that offline KL-regularized methods must overcome. To address this, ? proposes Weighted Preference Optimization (WPO), which enhances off-policy DPO by reweighting preference pairs based on their probability under the current policy, effectively simulating on-policy learning and improving the stability and efficiency of KL-regularized DPO. Furthermore, ? introduces Self-Play with

Adversarial Critic (SPAC), a provable and scalable offline alignment method for LLMs. SPAC leverages a Stackelberg game formulation with an adversarial critic and "on-average pessimism" to achieve theoretical convergence guarantees under weak data assumptions, while its practical DPO-like implementation bridges the gap between theoretical robustness and computational scalability for KL-constrained objectives.

Beyond the core optimization, fine-grained control and reward redistribution are crucial for effective KL-constrained RLHF. ? introduces Reinforced Token Optimization (RTO), which reframes RLHF as a Markov Decision Process (MDP) with token-wise rewards. It innovatively uses DPO to extract these fine-grained token-level signals for subsequent PPO training, addressing the sparsity of rewards while operating within the context of KL-regularized objectives. Similarly, ? proposes Attention Based Credit (ABC) and ? introduces RED (REward reDistribution), both aiming to generate dense, token-level reward signals from sparse, holistic feedback by leveraging the reward model's internal states or incremental score changes. These methods improve credit assignment and training stability in RLHF, providing a richer signal for the policy optimization that is often regularized by KL divergence.

While not explicitly KL-constrained, ?'s Contrastive Preference Learning (CPL) offers an alternative theoretical paradigm by directly learning policies from regret-based preferences without traditional RL. Its foundation in Maximum Entropy RL, which often involves entropy regularization, is conceptually related to the principles underlying KL-constrained objectives, demonstrating a broader theoretical exploration of preference modeling beyond direct reward maximization.

In conclusion, the theoretical foundations of KL-constrained preference learning have significantly advanced, moving from initial empirical successes to rigorous mathematical formulations and provable algorithms. Works like ? and ? provide the core theoretical understanding and guarantees for stability and efficiency. Subsequent research addresses critical challenges such as algorithmic bias ?, iterative learning efficiency ???, the online-offline gap ????, and fine-grained reward assignment ????. Despite these advancements, challenges remain in fully generalizing theoretical guarantees to the complex, non-linear

dynamics of all LLM architectures, dynamically adapting KL coefficients to varying tasks, and ensuring robustness to imperfections in reward models. Future work will likely focus on these areas to further solidify the theoretical robustness and practical efficacy of KL-constrained alignment.

5 Advanced Feedback Mechanisms and Robust Alignment

5.1 Reinforcement Learning from AI Feedback (RLAIF)

Reinforcement Learning from AI Feedback (RLAIF) marks a significant evolution in the alignment of large language models (LLMs), shifting the paradigm from reliance on costly and time-consuming human annotators to leveraging LLMs themselves for generating preference labels. This approach directly addresses the prohibitive costs and time associated with purely human-driven annotation, offering substantial scalability benefits for aligning ever-larger models [?](#). RLAIF aims to create more autonomous and cost-effective alignment pipelines, crucial for scaling with the rapid growth of LLMs.

The foundational work by lee2023mrw empirically demonstrated that RLAIF can achieve performance comparable to traditional Reinforcement Learning from Human Feedback (RLHF) across tasks such as summarization and dialogue generation. They introduced "direct RLAIF" (d-RLAIF), which streamlines the alignment process by having an LLM directly provide reward signals during the reinforcement learning phase, thereby circumventing the need for a separate reward model and mitigating the "reward model staleness" issue. To enhance the quality of AI-generated preferences, lee2023mrw explored techniques like Chain-of-Thought (CoT) reasoning and detailed preambles, aiming to maximize their alignment with human judgments. Theoretically, the efficacy of preference-based learning, which underpins both RLHF and RLAIF, can be framed as an Online Inverse Reinforcement Learning problem with known dynamics, providing a formal basis for understanding how models learn from comparative feedback [?](#).

Despite its promise, the transition to AI-generated feedback introduces a unique set of challenges, primarily concerning the nature and potential amplification of biases inherent in LLMs. saito2023zs7 critically identified and quantified "verbosity bias," demonstrating that AI models, when acting as labelers, often exhibit a strong preference for longer responses irrespective of their actual quality. This bias significantly diverges from human preferences and highlights a crucial limitation: unaddressed biases in AI feedback can lead to policies generating verbose and suboptimal outputs, compromising true alignment. Beyond explicit biases like verbosity, RLAIF faces the "verifier's dilemma": how to ensure the AI providing feedback is itself robustly aligned and free from subtle, systemic biases without an infinite regress of AI verifiers. The risk of self-reinforcing bias loops is particularly salient, where an AI labeler's idiosyncratic preferences or misalignments could be amplified in the policy model, which might then inadvertently influence the training of subsequent AI labelers, leading to a drift away from true human intent over generations of models. This necessitates robust mechanisms to verify the quality and impartiality of AI-generated feedback against an external, human-grounded standard.

To address these limitations and enhance the robustness and helpfulness of AI feedback, advanced methodologies have emerged. li2024ev4 introduced Hybrid Reinforcement Learning from AI Feedback (HRLAIF), a multi-stage framework designed to improve the accuracy of AI annotations, especially for complex judgments like factual correctness. HRLAIF employs a three-stage helpfulness labeling process that includes correctness verification against standard answers and reasoning process preference labeling, alongside an AI-driven red teaming approach for harmlessness. This structured approach significantly boosts AI's accuracy in critical categories like math and multiple-choice questions, mitigating the helpfulness degradation observed in basic RLAIF. Similarly, cao2024lh3 proposed RELC (Rewards from Language model Critique), which leverages LLMs' critique capabilities to generate dense, intermediate-step intrinsic rewards. This method directly addresses the sparse reward problem inherent in text generation, leading to more efficient and stable RL training by providing fine-grained feedback that would otherwise be prohibitively expensive or impossible to obtain from humans. These innovations demon-

strate a concerted effort to imbue AI feedback with greater reliability and granularity, moving beyond simple preference comparisons to more nuanced, verifiable assessments.

While RLAIF primarily focuses on leveraging AI for feedback generation, its integration into the broader RL-based alignment landscape also intersects with efforts to enhance efficiency. For instance, Proxy-RLHF by zhu2024zs2 offers an alternative avenue for computational efficiency by decoupling generation and alignment using a lightweight proxy model, which could complement RLAIF by providing a more efficient policy optimization step, even if the feedback itself is AI-generated. The versatility of RLAIF is also being explored in novel domains, such as autonomous driving, where multimodal AI feedback can be generated to align agents with complex environmental interactions and safety protocols ?. This showcases RLAIF’s potential to extend beyond text-only domains, enabling alignment in scenarios where human annotation is even more challenging or dangerous.

In conclusion, RLAIF represents a compelling advancement towards scalable and autonomous LLM alignment, offering a viable alternative to human-intensive RLHF by harnessing the generative and evaluative capabilities of LLMs themselves. Its core promise lies in its ability to overcome the cost and time bottlenecks of human annotation, as demonstrated by early successes in achieving comparable performance. However, this paradigm shift introduces critical challenges unique to AI-generated feedback, notably the risk of amplifying inherent LLM biases like verbosity and the fundamental "verifier’s dilemma" of ensuring the quality and impartiality of the AI labeler. While advanced methodologies like HRLAIF and RELC are actively addressing these issues by enhancing the accuracy and granularity of AI feedback, the field continues to grapple with preventing self-reinforcing misalignment loops and ensuring that AI-generated preferences truly reflect human values. Future research must focus on developing robust verification frameworks, understanding the long-term impact of AI-generated feedback on model capabilities and diversity, and designing mechanisms to prevent the subtle drift of alignment objectives when the feedback loop is predominantly AI-driven. This ongoing effort is crucial for realizing RLAIF’s full potential as a truly autonomous and trustworthy alignment pipeline.

5.2 Mitigating Reward Model Imperfections and Overoptimization

The initial successes of Reinforcement Learning from Human Feedback (RLHF) quickly unveiled a critical vulnerability: the tendency of large language models (LLMs) to exploit imperfections in proxy reward models, leading to "reward hacking" or "overoptimization" [1]. This phenomenon results in models generating misaligned or undesirable outputs despite achieving high scores on the proxy reward, thereby undermining the trustworthiness and reliability of aligned models.

Early investigations empirically demonstrated the prevalence of such issues. [2] rigorously showed that a significant portion of reported RLHF improvements could be attributed to models optimizing for spurious correlations, such as increased output length, rather than genuine quality. Their "Length-Only PPO (LPPO)" experiment strikingly reproduced most RLHF gains by simply optimizing for length, highlighting the non-robustness of reward models to shallow biases. Extending this diagnostic work, [3] revealed that reward overoptimization is not exclusive to traditional RLHF but also plagues Direct Alignment Algorithms (DAAs), which bypass explicit reward models. They established scaling laws for DAAs, demonstrating that performance degradation can occur rapidly, sometimes within a single training epoch, and that these models also exploit simple features like response length.

Addressing the root cause of reward model imperfections, [4] proposed enhancing reward model quality and interpretability through *self-generated critiques*. Their Critic-RM framework leverages an instruction-finetuned LLM to generate and filter high-quality critiques, then jointly fine-tunes the reward model on both scalar reward prediction and critique generation, leading to more robust and data-efficient reward signals.

Other strategies focus on regularizing the policy or constraining the optimization process to prevent models from venturing into unreliable regions of the reward landscape. For scenarios involving multiple objectives, [5] introduced constrained RLHF for composite reward models. Their method dynamically learns Lagrange multipliers to prevent individual reward components from overoptimizing beyond empirically identified "proxy

points," offering a more nuanced control than static weighting. ? tackled reward hacking through principled *reward shaping*, proposing "Preference As Reward (PAR)." This technique applies a sigmoid function to centered rewards, theoretically proven to stabilize critic training and minimize policy gradient variance, thereby mitigating the exploitation of reward function loopholes. Further, ? addressed over-optimization caused by reward model extrapolation errors in out-of-distribution (OOD) regions. Their Behavior-Supported Policy Optimization (BSPO) uses value regularization to restrict policy iteration to the in-distribution region of the reward model, penalizing OOD values without affecting in-distribution ones, and comes with theoretical convergence guarantees.

A prominent class of solutions involves quantifying the reward model's uncertainty to guide policy optimization away from unreliable regions. ? pioneered Uncertainty-Penalized RLHF (UP-RLHF), which augments the standard RLHF objective with an uncertainty regularization term. They introduced a novel Diverse Reward LoRA Ensemble to efficiently quantify reward model uncertainty, penalizing rewards based on the estimated standard deviation across the ensemble. Building on this, ? aimed to overcome the computational overhead of ensemble-based uncertainty quantification. They proposed a lightweight method to estimate reward uncertainty using *only the last layer embeddings* of a single reward model. This efficient uncertainty estimation then powers Adversarial Policy Optimization (AdvPO), a distributionally robust optimization framework that adversarially searches for the most pessimistic reward within a confidence region, proving to be less pessimistic than prior sample-wise penalization methods.

More recently, research has begun to explore internal, mechanistic approaches to combat reward hacking. ? identified the "Energy Loss Phenomenon" within the LLM's final layer as an internal signature of reward hacking, where excessive energy loss correlates with reduced contextual relevance. They proposed Energy loss-aware PPO (EPPO), which directly penalizes the increase in this internal energy loss during reward calculation. This novel approach moves beyond external reward signals or output-space regularization, offering a deeper, model-centric method to ensure that optimization genuinely aligns with human preferences.

Despite these advancements, the challenge of perfectly capturing and aligning with complex human preferences remains. While uncertainty quantification and constrained optimization improve robustness, they often rely on the quality of the underlying reward model or the definition of constraints. The novel internal mechanistic approaches, though promising, require further exploration into their generalizability and interpretability across diverse LLM architectures and tasks. Future work must continue to balance the efficiency of alignment algorithms with their robustness against subtle forms of reward exploitation, ensuring that LLMs are not only capable but also genuinely trustworthy and aligned.

5.3 Interpretable and Multi-Objective Reward Modeling

The foundational success of Reinforcement Learning from Human Feedback (RLHF) in aligning Large Language Models (LLMs) was initially built upon single-scalar, black-box reward models (RMs) ???. While effective for initial alignment, this approach suffered from significant limitations: it offered minimal insight into *why* a particular response was preferred, struggled to capture the inherent complexity and subjectivity of human values, and was prone to issues like reward hacking and the exploitation of spurious correlations, such as verbosity bias ???. Critically, this reduction of diverse human preferences to a single scalar often leads to an "objective mismatch," where the numerical optimization targets diverge from true human intent, resulting in unintended behaviors like excessive refusal or "laziness" ?. Furthermore, the opacity of these models raises concerns about whose values are encoded and the potential for unexamined biases ?. Consequently, research has increasingly focused on developing reward models that are not only accurate but also transparent, interpretable, and capable of handling multiple, potentially conflicting, human objectives (e.g., helpfulness, safety, conciseness) ???.

Addressing the inherent diversity and potential contradictions in human preferences has been a primary driver for multi-objective reward modeling. Early efforts moved beyond a monolithic reward by learning multiple reward functions or employing diverse evaluators. chakraborty20247ew introduced MaxMin-RLHF, which learns a mixture of reward models to represent distinct human sub-population preferences, thereby acknowl-

edging the insufficiency of a single reward for diverse feedback. Extending this concept, boldi2024d0s proposed Pareto-Optimal Preference Learning (POPL), which leverages lex-case selection to learn a set of Pareto-optimal reward functions or policies. This framework effectively handles contradictory preferences arising from "hidden context" without requiring explicit group labels, offering a more generalized approach to preference diversity. For complex tasks characterized by sparse and inconsistent human feedback, lai2024ifx introduced ALaRM (Align Language Models via Hierarchical Rewards Modeling), which integrates holistic rewards with proactively selected aspect-specific rewards, providing more precise and consistent guidance by decomposing the overall reward into finer-grained components. These methods collectively aim to represent the multifaceted nature of human preferences more faithfully than a single-scalar approach.

Beyond modeling diverse preferences, a critical challenge lies in balancing conflicting objectives, such as helpfulness and safety, and mitigating specific biases. tan2025lk0 addressed the "over-safe" phenomenon, where LLMs excessively refuse benign queries, through Equilibrate RLHF. This framework employs a Fine-grained Data-centric (FDC) approach to categorize safety data and an Adaptive Message-wise Alignment (AMA) method using gradient masking to selectively highlight safety-critical segments, allowing for nuanced control over safety responses. Complementing this, zhang2024b6u proposed Bi-Factorial Preference Optimization (BFPo), an efficient supervised learning framework that re-parameterizes the joint RLHF objective for safety and helpfulness into a single loss using a novel labeling function, significantly reducing the need for costly "red teaming" data while still managing multiple objectives. Similarly, xu20242yo developed Constrained Generative Policy Optimization (CGPO) with a "Mixture of Judges" (combining rule-based and LLM-based evaluators) to provide fine-grained control and enforce constraints, mitigating reward hacking in multi-task learning scenarios by integrating explicit objective-level feedback.

Mitigating specific biases, particularly verbosity, has also been a key focus. hou2024tvy incorporated "Bucket-Based Length Balancing" into their ChatGLM-RLHF pipeline, which involves grouping responses by length during reward model training to reduce the model's

inclination to prefer longer responses in a production setting. Further, chen2024vkb introduced length-regularized reward shaping within an iterative DPO framework, which directly debiases the implicit DPO reward model to construct a length-unbiased preference dataset, offering an alternative to loss-function-based length penalties and avoiding their associated hyperparameter tuning. These explicit adjustments to reward objectives are crucial for ensuring that high scores genuinely reflect quality rather than superficial characteristics.

The pursuit of truly interpretable and steerable reward models has also led to frameworks that provide granular insights into preference rationales. yu20249l0 enhanced reward model quality and interpretability with Critic-RM, which leverages self-generated, high-quality critiques to provide explicit rationales alongside scalar scores. This allows for a more transparent understanding of the reward model’s judgment process. In a similar vein, sun20238m7 introduced Principle-Driven Self-Alignment, where LLMs are guided by a small set of human-written principles (e.g., ethical, informative) and in-context exemplars during response generation. These principles are then "engraved" into the model’s parameters through fine-tuning, enabling the model to directly generate aligned responses without needing explicit prompting with principles during inference. This approach offers a direct path to multi-objective and steerable alignment by internalizing human-defined values. Furthermore, pignatelli2024ffp demonstrated the potential of LLMs for automated, granular credit assignment and subgoal decomposition, providing more interpretable and aspect-specific reward signals for complex tasks. This aligns with the idea of multi-task reward modeling, where hou202448j trained a unified reward model using a multi-task objective, combining pairwise ranking loss for human preferences with cross-entropy loss for binary-labeled reasoning data, including process supervision for intermediate steps. This integration of diverse feedback types contributes to a more comprehensive and interpretable reward signal.

A significant advancement in this direction is the comprehensive framework by wang20247pw, ArmoRM, which moves beyond black-box single-scalar RMs by training on multi-dimensional *absolute-rating* data. Each dimension corresponds to a human-interpretable objective

(e.g., helpfulness, correctness, verbosity, safety), providing decomposable explanations for reward scores and offering granular transparency into the model’s preference judgments. To enable dynamic and context-aware preference modeling, ArmoRM integrates a Mixture-of-Experts (MoE) Scalarization mechanism. This involves a shallow MLP gating network that dynamically selects and weights the most suitable reward objectives based on the input prompt’s context, allowing for flexible and steerable alignment. Crucially, wang20247pw also directly addressed the pervasive verbosity bias by explicitly adjusting each reward objective to ensure zero Spearman correlation with the verbosity objective, ensuring that high scores genuinely reflect quality rather than mere length. This framework achieved state-of-the-art performance on benchmarks like RewardBench, even surpassing much larger models and LLM-as-a-judge methods, demonstrating its efficiency and effectiveness in achieving nuanced and controllable alignment.

The evolution towards interpretable and multi-objective reward modeling represents a crucial step in developing more transparent, controllable, and robust LLMs that truly reflect the multifaceted nature of human preferences. While significant progress has been made in decomposing reward signals, dynamically weighting objectives, and mitigating specific biases, challenges remain. Future research needs to focus on the scalability of collecting high-quality, multi-dimensional absolute-rating data, developing more sophisticated theoretical guarantees for dynamic objective weighting, and establishing robust evaluation metrics for the interpretability and steerability of these complex reward models.

5.4 Implicit Human Signals and Distributionally Robust Alignment

Achieving robust and nuanced human-AI alignment necessitates moving beyond simplistic feedback mechanisms to capture subtle human preferences, while simultaneously ensuring model resilience against shifts in real-world data distributions. This subsection explores innovative approaches that integrate implicit human signals into reward models and leverage distributionally robust optimization techniques to enhance the generalization capabilities of aligned Large Language Models (LLMs).

To address the inherent limitations of explicit human feedback, such as high cost, scalability issues, and potential inconsistencies, researchers are exploring implicit behavioral cues. ? introduces *GazeReward*, a pioneering framework that directly integrates implicit eye-tracking (ET) data into the Reward Model (RM) for LLM alignment. A key innovation in this work is the use of ET prediction models to generate *synthetic* ET features from text. This synthetic generation effectively circumvents the practical challenges of collecting real human gaze data, making the integration of these subconscious preferences both scalable and cost-efficient, and demonstrating substantial performance gains in RM accuracy.

While richer feedback signals are crucial, the reward models themselves must be robust in interpreting these preferences. Traditional preference learning often assumes a linear scaling between preference strength and reward differences, an assumption that can lead to inflexible reward functions and misalignment between reward model accuracy and policy performance. ? addresses this by proposing an adaptive preference loss function, derived from a KL-constrained Distributionally Robust Optimization (DRO) formulation, which incorporates instance-specific scaling parameters. This allows the reward model to learn more nuanced relationships, assigning smaller scaling parameters to ambiguous preferences and larger ones to clear preferences, thereby enhancing the fidelity and internal robustness of the reward model’s interpretation of human feedback.

Building on the need for robustness, a significant challenge for aligned LLMs in real-world deployment is their susceptibility to out-of-distribution (OOD) prompt shifts, which can severely degrade performance. To counter this, ? introduces a comprehensive distributionally robust framework for Reinforcement Learning with Human Feedback (RLHF). This work applies DRO to the entire RLHF pipeline, including reward estimation, policy optimization, and Direct Preference Optimization (DPO), by defining uncertainty sets based on Total Variation (TV) distance. By explicitly accounting for shifts in the prompt distribution, this methodology aims to make aligned LLMs more resilient to unseen or OOD prompts, providing provable algorithms that enhance generalization capabilities with minimal alterations to existing pipelines.

Collectively, these advancements push the boundaries of human-AI alignment by leveraging richer, more subtle forms of feedback and ensuring robustness in diverse real-world scenarios. While ? pioneers the use of implicit human signals through synthetic generation to enrich the feedback loop, ? enhances the internal robustness of reward models by adaptively interpreting preference strengths. Complementing these, ? tackles external robustness, ensuring that the entire RLHF process and the resulting LLM generalize effectively to shifts in input data distributions. Despite these strides, challenges remain, particularly in bridging the fidelity gap between synthetic and real implicit signals, and in developing unified frameworks that seamlessly integrate implicit feedback with robust optimization techniques to address both internal preference ambiguity and external distribution shifts simultaneously. Further research is also needed to scale DRO techniques for even larger models and more complex, dynamic real-world environments.

6 RL for Enhanced LLM Capabilities and Specialized Applications

6.1 Enhancing LLM Reasoning with RL

Large Language Models (LLMs) excel in language generation, yet cultivating genuine, robust reasoning—encompassing mathematical problem-solving, multi-step logical deduction, and complex decision-making—presents a significant challenge for purely supervised learning paradigms. Reinforcement Learning (RL) has emerged as a powerful approach to address this limitation, enabling LLMs to ‘think’ more effectively and reliably by learning from diverse forms of feedback. Initial research, such as the development of the Chain-of-Thought Hub, empirically demonstrated a strong correlation between model scale and reasoning capabilities, and, crucially, highlighted the significant impact of Reinforcement Learning from Human Feedback (RLHF) on top-tier models like GPT-4 ?. This foundational understanding motivated deeper exploration into how RL could be specifically tailored to enhance reasoning.

A key avenue for improving LLM reasoning involves leveraging search and planning-based RL methods that allow models to explore complex solution spaces. havrilla2024m0y conducted a systematic comparative study of Expert Iteration (EI), Proximal Policy Optimization (PPO), and Return-Conditioned RL (RCRL) on mathematical word problems. Their findings revealed that Expert Iteration consistently outperformed PPO, particularly in deterministic reasoning environments, by effectively distilling high-quality reasoning trajectories. This work critically identified limited exploration during RL training as a significant bottleneck, suggesting that models often struggle to venture beyond solutions already produced by supervised fine-tuning (SFT) models, thereby limiting the full potential of RL. Building on the power of search, zhang2024q0e introduced LLaMA-Berry, a framework designed for Olympiad-level mathematical reasoning. It integrates Self-Refine Monte Carlo Tree Search (SR-MCTS) for robust exploration of the solution space, treating entire solutions as states and self-refinement as an optimization action. This is combined with a Pairwise Preference Reward Model (PPRM) trained via Direct Preference Optimization (DPO) and an Enhanced Borda Count for aggregating preferences, enabling smaller LLMs to tackle highly complex problems by efficiently navigating the vast search space of potential solutions.

For multi-step reasoning tasks, the granularity of feedback is paramount. Approaches that provide fine-grained, process-based feedback have shown particular promise. hao2025lc8 introduced RL-of-Thoughts (RLoT), a novel framework that models long-sequence reasoning as a Markov Decision Process (MDP). A lightweight "navigator model," trained with Double-Dueling-DQN, dynamically selects human cognition-inspired "logic blocks" (e.g., "Decompose," "Refine") based on the LLM's self-evaluation and a Process Reward Model. This inference-time guidance allows for adaptive, task-specific reasoning structures without modifying the underlying LLM, enabling smaller models to achieve performance comparable to much larger ones by dynamically adapting their reasoning strategy. Further advancing fine-grained feedback, chen20244ev proposed Step-level Value Preference Optimization (SVPO). This method leverages MCTS to *autonomously generate fine-grained step-level preferences* and explicitly integrates a value model into a novel DPO-like loss

function. By providing more granular feedback than coarse solution-level preferences, SVPO significantly boosts mathematical reasoning, demonstrating that learning from the intermediate steps of a reasoning chain is crucial for robust performance. In a similar vein, chen2025vp2 developed Solver-Informed Reinforcement Learning (SIRL) for optimization modeling. SIRL leverages external optimization solvers as objective verifiers, providing precise and comprehensive feedback signals—including syntax correctness, feasibility, and solution quality—as direct rewards. This application of Reinforcement Learning with Verifiable Rewards (RLVR) provides highly granular, objective feedback at each step of the model generation process, grounding LLMs for authentic optimization modeling and significantly improving their functional correctness.

Beyond external feedback, methods that enable LLMs to self-verify and learn from internal signals are critical for robust reasoning. liu2025lxxv introduced RISE (Reinforcing Reasoning with Self-Verification), a novel online RL framework that explicitly and simultaneously trains an LLM to improve both its problem-solving and self-verification abilities. RISE leverages verifiable rewards not only to guide solution generation but also to align the model's self-verification ability on-the-fly, addressing the challenge of "superficial self-reflection" where LLMs struggle to robustly identify flaws in their own reasoning. However, the utility of purely internal feedback requires careful consideration. zhang2025d44 critically examined Reinforcement Learning from Internal Feedback (RLIF), exploring self-certainty and entropy-based signals as reward proxies. Their theoretical and empirical analysis revealed that these internal signals primarily drive the model to minimize policy entropy, making the output distribution more deterministic. While RLIF can initially boost reasoning performance for *base LLMs*, this gain often degrades with continued training, and it shows little to no improvement for instruction-tuned models, highlighting a "no free lunch" scenario where the effectiveness of internal feedback is highly dependent on the model's initial state and training dynamics.

RL has also proven effective in enhancing reasoning in specific domains and necessitates specialized data infrastructure. anand2024rnl applied Reinforcement Learning with Human and Artificial Intelligence Feedback (RLHAIF) using PPO, DPO, and Remax op-

timization to improve LLM performance on physics problem-solving. Their MISTRAL-PPO model demonstrated marked improvements in reasoning and accuracy, showcasing RL's potential in complex STEM domains. To support such specialized applications, albalak2025wyc addressed the critical need for high-quality, task-specific data by creating Big-Math, the largest open-source dataset of over 250,000 math problems specifically curated for RL training. This rigorous human-in-the-loop filtering and novel reformulation of multiple-choice questions into open-ended formats bridge the quality-quantity gap, providing essential infrastructure for robust reasoning research.

In summary, the application of RL to LLM reasoning has evolved from empirical observations of RLHF's importance to sophisticated algorithmic interventions. These include search-based methods for robust exploration, fine-grained and process-based feedback mechanisms for multi-step tasks, and frameworks for self-verification and internal feedback. While significant progress has been made in enhancing mathematical, logical, and multi-step reasoning, critical challenges remain. The limited exploration identified by havrilla2024m0y persists, particularly in complex, open-ended reasoning environments. Furthermore, a crucial critical perspective from bao2024wnc suggests that post-training techniques like SFT and RLHF may inadvertently *weaken* the ideal causal structure of reasoning, potentially promoting "explaining" over genuine "reasoning." This highlights the need for future work to focus not only on improving performance but also on ensuring that RL-driven enhancements foster truly robust, causally sound, and interpretable reasoning processes in LLMs, rather than merely optimizing for superficial correctness. Future research will likely explore more sophisticated exploration techniques, integrate explicit value learning more robustly, and develop methods that explicitly preserve or strengthen the causal integrity of LLM reasoning.

6.2 RL for Tool Use and Agentic Behavior

Reinforcement learning (RL) is pivotal in transforming Large Language Models (LLMs) from static text generators into dynamic, interactive agents capable of effectively using external tools and exhibiting sophisticated agentic behaviors. This involves learning to

orchestrate diverse tools, stabilizing multi-turn interactions, and adaptively planning actions in complex environments, often optimizing for conversation-level preferences.

A foundational challenge in building truly interactive and goal-oriented AI agents is managing multi-turn interactions and aligning them with human preferences over extended dialogues. ? addresses this by introducing Multi-turn Preference Optimization (MTPO), a novel policy optimization algorithm for multi-turn interactions. This work formalizes multi-turn preference-based RL within a Contextual Markov Decision Process (CMDP) framework, crucially utilizing *conversation-level* preference feedback rather than single-turn feedback, and proposing a novel preference-based Q-function that considers long-term consequences. MTPO provides the first theoretically grounded policy optimization algorithm with convergence guarantees for general multi-turn preference-based RL, enabling models to stabilize complex sequential interactions.

Building on the ability to manage multi-turn interactions, RL further empowers LLMs to interact with and orchestrate external tools. In the visual domain, ? presents OpenThinkIMG, a comprehensive framework that enables Large Vision-Language Models (LVLMs) to dynamically and adaptively utilize external visual tools. Their V-TOOLRL framework, which employs Group-wise Proximal Policy Optimization (GRPO), allows LVLMs to learn adaptive tool invocation policies, moving beyond the limitations of static supervised fine-tuning (SFT) by optimizing directly for task success using feedback. This approach significantly enhances an agent's ability to "think with images" and orchestrate diverse visual tools for complex problem-solving. While ? focuses on orchestrating existing tools, the challenge remains for models to generate and refine complex outputs through iterative tool use.

Addressing this, ? introduces the Reasoning-Rendering-Visual-Feedback (RRVF) framework, which enables Multimodal Large Language Models (MLLMs) to learn complex generative tasks, such as image-to-code generation, solely from raw images without relying on expensive image-text supervision. RRVF leverages the "Asymmetry of Verification" principle, where verifying a rendered output against an image is easier than generating the code from scratch. The framework employs a closed-loop iterative process where the MLLM generates code, external tools render it, and a "Visual Judge" MLLM pro-

vides natural language feedback, all optimized end-to-end using GRPO with a hybrid reward function. This demonstrates a sophisticated form of agentic behavior through self-correction and iterative refinement, allowing models to adaptively plan actions to achieve a desired visual outcome. Both **?** and **?** highlight the effectiveness of RL in enabling tool use, but they also implicitly rely on powerful external LLMs (like GPT-4o) for either initial action planning or reward signal generation, pointing to a bootstrapping dependency.

Beyond mere tool invocation, RL also facilitates LLMs in exhibiting more complex, strategic agentic behaviors, including interactions with other systems. **?** explores this by presenting **CheatAgent**, a novel framework where LLMs act as intelligent, black-box adversarial agents to attack LLM-empowered recommender systems. Unlike traditional RL agents that struggle with complex textual inputs, **CheatAgent** leverages LLMs' inherent language comprehension and open-world knowledge. It optimizes an auxiliary LLM's "policy" (represented by a trainable prefix prompt) through self-reflection based on the victim system's loss, akin to policy gradient methods. This work demonstrates a distinct form of agentic behavior where LLMs strategically interact with their environment and other systems, moving beyond predefined tool use to exhibit adaptive planning for specific goals, even adversarial ones.

In conclusion, RL is instrumental in advancing LLMs beyond static text generation towards dynamic, interactive, and goal-oriented AI agents. Significant progress has been made in stabilizing multi-turn interactions and optimizing for conversation-level preferences through methods like MTPO **?**. Furthermore, RL has enabled sophisticated tool orchestration, as seen in visual domains with **OpenThinkIMG** **?** and iterative self-correction for generative tasks with RRVF **?**. The emergence of LLMs as strategic agents, exemplified by **CheatAgent** **?**, underscores the breadth of agentic behaviors RL can foster. Unresolved challenges include bridging theoretical guarantees with practical deep RL implementations, reducing reliance on powerful external models for feedback or data generation, and further exploring the ethical implications of increasingly autonomous and strategic AI agents.

6.3 LLM/VLM-Driven Automated Reward Design

The manual design of reward functions remains a significant bottleneck in Reinforcement Learning (RL), often requiring extensive human effort, domain expertise, and iterative refinement to avoid suboptimal or unintended agent behaviors. Recent advancements in large language models (LLMs) and vision-language models (VLMs) offer promising avenues to automate this challenging process, bridging the gap between high-level human intent and low-level reward signals for autonomous agents.

Early efforts in this domain leveraged LLMs' code generation capabilities to synthesize executable reward functions. ma2023vyo introduced *Eureka*, a pioneering method where an LLM (GPT-4) generates executable Python reward code by taking environment source code as context. This approach employs an evolutionary search and a novel "reward reflection" mechanism, which provides fine-grained feedback on policy dynamics to the LLM for targeted code editing, enabling human-level and even superhuman performance on complex dexterous manipulation tasks. Building upon the concept of LLM-driven reward evolution, hazra2024wjp proposed *REvolve*, an evolutionary framework that uses LLMs as intelligent genetic operators for mutation, crossover, and selection of reward functions. Critically, *REvolve* addresses a limitation of earlier iterative methods by employing human feedback (preferences over policy rollouts) as a direct, non-differentiable fitness function, making it particularly effective for tasks with subjective notions of "good" behavior and reducing the need for a pre-defined fitness metric.

A parallel and distinct line of research shifted from LLMs generating explicit code to VLMs directly inferring reward signals from visual observations, thereby eliminating the need for environment source code or low-level state access. rocamonde2023o9z demonstrated that pretrained VLMs, such as CLIP, can serve as *zero-shot reward models (VLM-RMs)* for vision-based RL. Their method computes reward as the cosine similarity between a natural language task description's embedding and the visual observation's embedding, introducing "Goal-Baseline Regularization" to project out irrelevant information and shape the reward landscape more effectively. This approach highlighted a strong scaling effect, where larger VLMs yielded superior reward models for complex visual tasks.

However, direct VLM scores can sometimes be noisy or high-variance, prompting further innovation.

Addressing these limitations, wang2024n8c introduced *RL-VLM-F*, which leverages advanced VLMs (e.g., GPT-4V, Gemini) to provide automated *preference feedback* over pairs of visual observations based on a natural language goal. Instead of directly outputting reward scores, the VLM acts as an automated human annotator, generating preference labels from which a robust reward function is then learned. This method employs a novel two-stage VLM querying process for analysis and labeling, operating solely on visual observations and a text goal, making it highly applicable to complex visual tasks, including deformable object manipulation, where ground-truth state information is often unavailable. *RL-VLM-F* significantly reduces human effort in preference-based RL and demonstrates superior performance over prior methods that directly use VLM scores for reward generation.

Collectively, these works represent a significant leap in automating reward design, moving from LLMs generating interpretable code that requires environment state access to VLMs directly interpreting visual observations and natural language goals. While LLM-driven code generation offers transparency and fine-grained control over reward components, VLM-driven inference provides a more natural, visual-centric approach, particularly for tasks where only pixel data is available. Future research could explore hybrid approaches that combine the interpretability of LLM-generated code with the visual grounding of VLM-inferred signals, or focus on improving the robustness of VLM-based rewards to abstract visual environments and enhancing their spatial reasoning capabilities for real-world deployment. The computational cost associated with iterative LLM interactions and extensive RL training runs, as well as ensuring the safety and alignment of autonomously designed rewards, remain critical areas for continued investigation.

6.4 Domain-Specific Applications: Software Engineering and Code Generation

Large Language Models (LLMs) have demonstrated impressive capabilities in generating code, yet bridging the gap between syntactically plausible outputs and functionally correct, domain-specific performance in software engineering tasks remains a significant challenge. Reinforcement Learning (RL) emerges as a powerful paradigm to address this, enabling LLMs to learn from external feedback mechanisms, ranging from lightweight rule-based rewards to rigorous external verification, thereby enhancing their ability to fix bugs, resolve issues, and generate verifiable code.

A pioneering effort in applying RL to real-world software engineering challenges is presented by ? with their SWE-RL framework. This work tackles the impracticality of execution-based rewards for complex bug-fixing and issue resolution tasks by introducing a scalable, *lightweight, rule-based reward function*. Leveraging a massive, curated dataset of GitHub Pull Requests, SWE-RL employs `difflib.SequenceMatcher` to calculate patch similarity, providing an efficient proxy for correctness. This approach not only achieves state-of-the-art performance for medium-sized LLMs on the SWE-bench Verified benchmark (41.0

Building upon the potential of RL to imbue LLMs with domain-specific correctness, ? focuses on the highly specialized domain of Hardware Description Language (HDL) generation, specifically Verilog. This research addresses the stringent requirement for *functional correctness* by integrating external verification feedback directly into the RL training loop. Their methodology combines Supervised Fine-Tuning (SFT) with Direct Preference Optimization (DPO), where preference pairs are derived from a novel automatic testbench generation pipeline that incorporates real Verilog Compiler Simulator (VCS) feedback. This direct integration of verification insights allows the LLM to learn what constitutes functionally correct code, moving beyond the proxy rewards seen in approaches like SWE-RL. By leveraging DPO, ? effectively mitigates the "reward hacking" problem often associated with explicit reward models, demonstrating a robust framework

for aligning LLMs with verifiable functional performance. Nevertheless, the computational overhead of external simulators and the inherent limitation that automatically generated testbenches cannot guarantee coverage of all functional cases remain challenges.

These works collectively illustrate a crucial intellectual trajectory in reinforcement learning for language processing: the strategic application of RL to complex, real-world, domain-specific tasks where traditional, costly execution-based rewards are often infeasible. ? highlights the viability of scalable, lightweight, and rule-based reward functions as a practical alternative for driving LLM improvement in challenging environments like open-source software evolution, even fostering emergent generalized reasoning. Complementing this, ? showcases how RL, particularly DPO, can effectively integrate non-textual, domain-specific external feedback, like hardware simulation results, to achieve stringent functional correctness, a critical aspect beyond purely linguistic objectives. The unresolved tension in this domain lies in balancing the practicality and scalability of proxy reward functions with the high fidelity and computational intensity of true external verification, and in developing hybrid approaches that can leverage both to achieve robust, verifiable, and efficient domain-specific LLM performance. Future research will likely explore more efficient verification techniques, adaptive reward mechanisms, and multi-modal feedback integration to further bridge the gap between LLM generative power and the exacting demands of real-world engineering.

7 Efficiency, Scalability, and Responsible AI in RL for Language Processing

7.1 Training Efficiency and Parameter-Efficient RLHF

The application of Reinforcement Learning from Human Feedback (RLHF) has proven instrumental in aligning large language models (LLMs) with human preferences, yet its substantial computational cost and memory footprint pose significant barriers to broader accessibility and scalability ?. Addressing this, research has focused on three primary di-

rections: integrating parameter-efficient fine-tuning (PEFT) methods, developing efficient RLHF frameworks, and exploring alternative paradigms that simplify the RLHF pipeline.

A direct approach to mitigating the resource demands of RLHF involves the integration of Parameter-Efficient Fine-Tuning (PEFT) methods. ? introduced Parameter Efficient Reinforcement Learning from Human Feedback (PE-RLHF), demonstrating that Low-Rank Adaptation (LoRA) can be effectively applied to both reward model (RM) training and the subsequent reinforcement learning of the policy model. This innovation significantly reduces memory usage (20-57

Beyond parameter efficiency, advancements in RLHF frameworks and system optimizations are crucial for handling the complex, distributed nature of the training process. ? proposed Efficient Sampling-based RL (ESRL) for sequence generation, which is highly relevant to RLHF. ESRL employs a two-stage sampling process to decouple sequence sampling from gradient computation and dynamic sampling to adaptively adjust sampling size and temperature, leading to substantial reductions in memory consumption and training time. Building on the need for robust system architectures, ? introduced HybridFlow, a flexible and efficient RLHF framework. HybridFlow leverages a hierarchical hybrid programming model and a specialized 3D-HybridEngine to manage heterogeneous workloads and complex parallelism strategies, achieving throughput improvements of 1.53x to 20.57x over existing baselines. Further optimizing resource utilization, ? presented ReaL, a system for efficient RLHF training that dynamically reallocates LLM parameters across the training cluster. ReaL uses an MCMC-based search algorithm to discover optimal execution plans for fine-grained resource allocation and parallelization, yielding speedups of up to 3.58x and significantly improving performance over static heuristic approaches.

Another promising avenue for efficiency lies in simplifying the RLHF pipeline itself by bypassing the computationally intensive reward model inference step. ? explored zeroth-order optimization for direct policy learning from human ranking oracles, introducing ZO-RankSGD. This method directly uses human preferences to estimate descent directions for non-convex functions, offering a provably convergent approach without an

explicit reward model. Extending this, ? proposed Zeroth-Order Policy Gradient (ZPG) and Zeroth-Order Block-Coordinate Policy Gradient (ZBCPG) for general RL problems. These algorithms estimate policy gradients from human preferences over trajectory pairs, providing provably efficient policy-based RLHF without reward inference, even in stochastic environments. Complementing this, ? developed Batched Sequential Action Dueling (BSAD), a model-free algorithm that directly identifies the optimal policy from human preference information in episodic MDPs. BSAD achieves instance-dependent sample complexity comparable to classic RL by employing backward action dueling and reward-free exploration, further demonstrating the feasibility of reward-model-free RLHF.

Collectively, these advancements represent a significant stride towards democratizing RLHF. While PEFT methods like LoRA offer immediate memory and speed benefits, sophisticated frameworks such as HybridFlow and ReaL tackle the systemic challenges of distributed training. Simultaneously, direct policy optimization methods provide alternative, more streamlined pipelines by eliminating the reward model. Future research will likely focus on integrating these diverse efficiency strategies, exploring more advanced PEFT and Representation Fine-Tuning (ReFT) techniques, and developing adaptive systems that can dynamically choose the most efficient combination of methods based on model size, task complexity, and available hardware.

7.2 Dynamic and Inference-Time Alignment

The prevailing paradigm for aligning large language models (LLMs) with human preferences and values, predominantly through methods like Reinforcement Learning from Human Feedback (RLHF), typically results in static models. Once aligned, these models exhibit inflexibility to diverse user needs, evolving contexts, or real-time behavioral adjustments, necessitating costly and time-consuming retraining for any modification. This inherent inflexibility presents a significant operational and scalability challenge, especially in rapidly changing deployment environments. Consequently, an emerging and critical research direction focuses on enabling dynamic adjustment and control of LLM behavior directly during inference. These innovative inference-time alignment strategies offer un-

paralleled flexibility, efficiency, and real-time adaptability without the need for extensive post-training fine-tuning, marking a significant shift from static, pre-computed alignment to adaptive, context-aware control.

These dynamic alignment approaches can be broadly categorized based on their mechanism: direct inference-time interventions that manipulate model outputs or internal states, and inference-time search and self-optimization strategies that leverage the LLM’s own capabilities or auxiliary agents to refine behavior. Each category offers distinct advantages and challenges in dynamically controlling the policy learned, often, through RLHF.

7.2.1 Direct Inference-Time Interventions

This category encompasses methods that directly modify the model’s decoding process or internal activations to steer its behavior towards desired alignment objectives. These interventions are typically lightweight, operate on the fly, and often serve to dynamically adjust or correct the behavior of an already RLHF-aligned policy.

A foundational approach in this area is *Representation Engineering* or *Control Vectors*, which directly manipulate the internal activation space of LLMs to elicit specific behaviors **???**. These methods identify "steering vectors" by contrasting activations from desired and undesired responses (e.g., helpful vs. unhelpful, safe vs. unsafe). By adding or subtracting these vectors to the hidden states of an LLM during inference, the model’s subsequent generation can be steered towards the desired attribute. For instance, **?** demonstrated that simple inference-time interventions on internal states can significantly improve truthfulness in LLMs. While powerful for fine-grained control, a critical challenge lies in the interpretability of these steering vectors and the potential for unintended side effects or over-steering, as their impact on complex, emergent behaviors is not always fully understood. Moreover, the effectiveness often depends on careful calibration and the quality of the contrasting examples used to derive the vectors. These techniques offer a direct way to modify the latent policy of an RLHF-aligned model without altering its weights.

Decoding-time Realignment (DeRa) **?** provides a logit-level intervention for dynamic

adjustment of alignment strength during inference. Grounded in the theoretical insight that models aligned with varying KL regularization strengths are effectively geometric mixtures of a reference model and a single aligned model, DeRa linearly combines the logits of these two models during decoding. This enables real-time exploration of the reward-regularization trade-off, allowing users to dynamically dial up or down the "alignedness" of a model's output (e.g., helpfulness, conciseness) without retraining. While offering fine-grained control over an RLHF-trained policy, DeRa incurs increased inference time and memory usage due to the need to compute logits from two models simultaneously. Its effectiveness also relies on the assumption that a linear combination of logits accurately captures the desired behavioral spectrum, which might not hold universally across all tasks or models.

Addressing the critical aspect of safety, *InferAligner* ? introduced an inference-time solution for enhancing harmlessness through cross-model guidance. This method extracts "Safety Steering Vectors" (SSVs) from an already safety-aligned model (often trained with RLHF) and conditionally applies them to the target model's activations. InferAligner employs a "guidance gate" that activates this intervention only when harmful intent is detected in the input, preventing unnecessary shifts that could degrade performance on benign tasks. This conditional, activation-level steering effectively transfers safety knowledge without retraining the target model, providing a robust defense against "model hacking" attacks like Probability Manipulation (ProMan) ? that can bypass static alignment efforts. A key limitation is its dependency on the quality and robustness of the "teacher" safety-aligned model and the accuracy of the harmful intent detection mechanism. Poor detection could lead to negative transfer or missed harmful outputs. Unlike DeRa which operates on output logits for general alignment, InferAligner focuses on internal activations for a specific safety objective, demonstrating different granularities and targets of intervention.

7.2.2 Inference-Time Search and Self-Optimization

This category focuses on methods where the LLM itself, or a lightweight auxiliary model, dynamically optimizes its prompts, reasoning paths, or objectives during inference to achieve better alignment or task performance. These approaches often involve an iterative search process and can be seen as alternatives or complements to traditional RLHF.

Dynamic Rewarding with Prompt Optimization (DRPO) ? exemplifies a tuning-free, search-based framework for self-alignment at inference time. DRPO allows LLMs to iteratively self-improve by crafting optimal alignment instructions, including system prompts and in-context learning (ICL) examples, without any additional training or human intervention. Its core innovation lies in a dynamic rewarding mechanism that adjusts LLM-based rewards based on specific queries, identifying and rectifying model-specific alignment weaknesses. This enables LLMs to adapt efficiently to diverse alignment challenges, demonstrating that base models aligned with DRPO can achieve higher average alignment scores than their SFT/RLHF-tuned counterparts, challenging the necessity of extensive fine-tuning for alignment. However, a significant critical point is the potential for self-deception or bias amplification, as the LLM acts as both the generator and the rewarder. The iterative search process can also incur substantial computational cost at inference time, limiting its real-time applicability for latency-sensitive applications.

Beyond general alignment, dynamic control is also being applied to guide complex reasoning processes. *RL-of-Thoughts (RLoT)* ? is an inference-time technique that trains a lightweight "navigator model" using reinforcement learning to dynamically generate task-adaptive logical structures for LLM reasoning. By selecting human cognition-inspired "logic blocks" (e.g., "Decompose," "Debate," "Refine") based on the LLM's self-evaluation of its current reasoning status, RLoT enables adaptive and task-specific reasoning without modifying the base LLM. This allows smaller models to achieve reasoning performance comparable to much larger ones, showcasing the power of dynamic, RL-guided search in the thought space. The explicit use of RL to train the navigator model directly connects this method to the review's core theme. A limitation of RLoT is the complexity of designing effective "logic blocks" and the reliance on the LLM's self-evaluation accuracy, which

can be imperfect. Unlike DRPO which optimizes the *input instructions*, RLoT optimizes the *internal reasoning process* itself, demonstrating different facets of inference-time self-optimization.

In conclusion, the shift towards dynamic and inference-time alignment marks a significant evolution in LLM deployment, moving beyond the inflexibility of static post-training alignment. These approaches, ranging from real-time adjustment of alignment strength via logit manipulation [1] and cross-model safety guidance through activation steering [2], to autonomous self-improvement via prompt optimization [3] and dynamic reasoning navigation [4], offer unparalleled flexibility and efficiency. They address the core limitation of static alignment by allowing LLMs to adapt to diverse user needs and contexts in real-time, significantly reducing the operational overhead associated with maintaining multiple aligned models. Future research will likely focus on combining these diverse techniques, further reducing inference overhead, improving the robustness of self-evaluation and steering mechanisms, and broadening their applicability to a wider array of alignment objectives and complex, multi-modal scenarios, ultimately leading to more adaptable and context-aware AI systems.

7.3 Privacy, Security, and Robustness of Aligned LLMs

Ensuring the privacy, security, and robustness of large language models (LLMs), particularly those aligned with reinforcement learning from human feedback (RLHF), presents critical challenges for their trustworthy real-world deployment. This section delves into frameworks for privacy-preserving alignment, investigations into data memorization, and the identification of vulnerabilities that can circumvent alignment mechanisms.

Addressing the inherent privacy risks in RLHF, [1] proposes a novel, end-to-end Differentially Private (DP) framework for aligning LLMs. This framework integrates DP across all three critical stages of RLHF: supervised fine-tuning (SFT), reward model training, and policy optimization via a modified Proximal Policy Optimization (DPPPO). A key technical contribution is the demonstration that the reward model itself must be trained with DP to ensure the overall privacy guarantees of the alignment pipeline, a crucial in-

sight for developing truly privacy-preserving LLMs. While providing strong mathematical privacy guarantees, the framework simplifies privacy accounting by assuming disjoint user contributions, leaving more complex multi-stage scenarios for future exploration.

Building upon the general concern for privacy, ? empirically investigates a specific privacy risk: memorization of training data within the RLHF pipeline, particularly for code completion tasks. Using a k -approximate counterfactual memorization definition, their work reveals that examples memorized during initial fine-tuning are likely to persist, but data used solely for reward model training is less prone to memorization by the final aligned model. This offers valuable guidance for mitigating privacy risks when handling sensitive user preferences, and critically, they find that direct preference optimization methods like IPO tend to increase memorization risk compared to RLHF.

Beyond privacy, the security and robustness of aligned LLMs are paramount. Efforts to make alignment more efficient, such as the Principle-Driven Self-Alignment proposed by ?, aim to instill helpful and ethical behaviors with minimal human supervision. Their SELF-ALIGN pipeline, particularly **Principle-Driven Self-Alignment** and **Principle Engraving**, allows LLMs to internalize abstract principles and generate aligned responses from scratch, addressing the scalability and cost limitations of traditional RLHF. However, the very robustness of such alignment, regardless of its efficiency, is a significant concern.

The efficacy of alignment mechanisms against sophisticated attacks is critically examined by ?, who demonstrate that even aligned open-sourced LLMs remain vulnerable to misuse. They introduce Probability Manipulation (ProMan), a novel "model hacking" attack that, with white-box access, directly manipulates internal logit probabilities during token generation. This allows attackers to force models to produce harmful or private content, fundamentally questioning the sufficiency of current alignment strategies (including RLHF and SFT) against such potent, internal attacks.

Further highlighting the fragility of alignment, ? reveals a novel and widespread vulnerability: non-standard Unicode characters can significantly reduce the efficacy of RLHF-implemented guardrails. Their empirical study across 15 diverse LLMs demonstrates that such inputs can lead to prompt leakage, hallucinations, and comprehension

errors, even exposing sensitive system prompts. This work provides strong evidence that current RLHF-based safety mechanisms are insufficient against subtle input perturbations, underscoring a critical gap in true language comprehension and adversarial robustness.

In conclusion, the literature reveals a complex interplay between privacy, security, and robustness in aligned LLMs. While frameworks like ? offer pathways for privacy-preserving alignment and ? provides insights into memorization dynamics, the works by ? and ? expose fundamental vulnerabilities that can circumvent even well-intentioned alignment efforts. An unresolved tension remains in developing alignment techniques that are not only efficient and private but also inherently resilient to both internal model manipulation and novel adversarial inputs, necessitating a paradigm shift towards more robust and adversarially-aware RLHF training and defense mechanisms.

7.4 Critical Evaluation and Sociotechnical Limits of Alignment

Despite the remarkable successes of Reinforcement Learning from Human Feedback (RLHF) in aligning Large Language Models (LLMs) with human preferences, a growing body of literature reveals inherent limitations, biases, and unintended consequences that challenge the very foundations of current alignment paradigms. This critical examination extends beyond purely technical performance, delving into the broader sociotechnical implications of AI safety.

One significant limitation of RLHF is the trade-off between generalization and diversity. kirk20230it conducted a systematic analysis, demonstrating that while RLHF significantly improves out-of-distribution (OOD) generalization, it simultaneously leads to a substantial reduction in output diversity. Building on this observation, mohammadi20241pk provided a mechanistic explanation for this loss of diversity, specifically focusing on creativity. Their work empirically showed that aligned models gravitate towards distinct, limited "attractor states" in their output embedding space, exhibiting behavior akin to mode collapse and transforming LLMs into more deterministic algorithms, thereby sacrificing creative potential for consistency.

Beyond these inherent trade-offs, RLHF is susceptible to various biases and vulnera-

bilities. saito2023zs7 identified a significant "verbosity bias" where LLMs, when used as evaluators in Reinforcement Learning from AI Feedback (RLAIF), tend to prefer longer, more verbose answers even if their quality is similar to shorter ones, highlighting a discrepancy with human preferences for conciseness. Furthermore, the integrity of RLHF can be compromised by malicious attacks. baumgrtner2024gu4 demonstrated that RLHF is highly vulnerable to stealthy "preference poisoning" attacks, where a small fraction of naturalistic, poisoned preference data can manipulate the model to generate specific entities with desired sentiments, posing significant security risks.

A fundamental conceptual challenge underlying many of these issues is the "objective mismatch" problem. lambert2023c8q formalized this problem, arguing that the decoupling of numerical objectives across reward model training, policy optimization, and downstream evaluation leads to unintended behaviors such as excessive refusals, "laziness," or verbosity. This mismatch arises from erroneous assumptions about the correlation between proxy rewards and true human intent. Reinforcing this, lambert2023bty provided a critical historical and conceptual analysis of RLHF reward models, highlighting the opacity of their value encoding process and the risks of applying optimization stacks designed for clear control problems to the vague domain of human language and values. They argue that the foundational assumptions—that human preferences are quantifiable and that reward maximization leads to desired behaviors—are often ill-posed in complex ethical contexts.

The impact of RLHF can also extend to the very internal reasoning mechanisms of LLMs. Surprisingly, bao2024wnc employed a causal inference framework to diagnose LLM reasoning with Chain-of-Thought (CoT) and found that while in-context learning strengthens ideal causal reasoning structures, post-training methods like Supervised Fine-Tuning (SFT) and RLHF can actually *weaken* them. This suggests a potential trade-off where optimizing for external metrics through RLHF might inadvertently degrade the underlying causal fidelity indicative of genuine reasoning, challenging the assumption of universal benefits from alignment.

Ultimately, these technical and conceptual limitations point to broader sociotechnical

limits of AI alignment. lindstrm20253o2 offers a comprehensive critique, questioning the sufficiency of the widely adopted "helpful, harmless, honest" (HHH) principles. They argue that the operationalization of HHH often oversimplifies complex ethical considerations and can even tolerate harm. By applying "the curse of flexibility" from system safety literature, they explain that the generalist nature and immense power of LLMs make it inherently difficult to define and ensure safety requirements through purely model-centric technical design. This work advocates for a paradigm shift, proposing that AI safety must be approached as a comprehensive sociotechnical discipline, integrating ethical, institutional, and process considerations alongside purely technical design, acknowledging that alignment is not solely a technical problem solvable within algorithmic boundaries.

In conclusion, while RLHF has been instrumental in advancing LLM capabilities, a critical evaluation reveals fundamental trade-offs between generalization and diversity, a loss of creativity, susceptibility to biases and attacks, and a pervasive "objective mismatch" problem. Furthermore, RLHF may inadvertently compromise the internal reasoning fidelity of models. These issues, coupled with the inherent vagueness of ethical principles, underscore that AI alignment is a complex sociotechnical challenge that cannot be resolved through technical design alone, necessitating a multidisciplinary approach that integrates broader ethical, institutional, and process considerations.

8 Conclusion and Future Directions

8.1 Summary of Key Advancements and Intellectual Trajectories

The application of reinforcement learning (RL) to language processing has undergone a profound transformation, evolving from early attempts at structured prediction and direct generation to becoming an indispensable tool for aligning and enhancing large language models (LLMs). This trajectory marks a continuous drive towards more capable, aligned, and versatile language AI, fundamentally reshaping the field.

Initially, the integration of RL for language generation faced challenges related to sparse rewards and the vast, discrete action spaces of language. The advent of Rein-

forcement Learning from Human Feedback (RLHF) marked a pivotal shift, enabling the alignment of LLMs with complex human preferences. Early RLHF implementations, often relying on Proximal Policy Optimization (PPO), proved effective but were notoriously complex, unstable, and sensitive to hyperparameters [1]. To address these limitations, significant algorithmic refinements emerged. Direct Preference Optimization (DPO) [2] simplified the RLHF pipeline by reformulating it as a classification loss, directly optimizing the policy without an explicit reward model or reinforcement learning. This innovation offered a more stable and computationally lightweight alternative to PPO, achieving comparable or superior performance. Building on this, [3] provided a rigorous theoretical analysis of KL-constrained RLHF, leading to iterative DPO and multi-step rejection sampling strategies that enhance exploration and robustness.

Beyond algorithmic simplification, efforts focused on improving the efficiency and scalability of RLHF. Reinforced Self-Training (ReST) [4] decoupled data generation and policy improvement into iterative phases, allowing for compute- and sample-efficient alignment by progressively refining policies on increasingly high-quality, reward-filtered data. Complementing these algorithmic advancements, infrastructural innovations were crucial. [5] introduced HybridFlow, a flexible and efficient framework that combines single-controller and multi-controller paradigms to manage the complex, distributed computation inherent in RLHF, significantly boosting throughput and simplifying development.

The reliance on extensive human feedback, a bottleneck for scalability, spurred innovations in feedback mechanisms and reward modeling. Reinforcement Learning from AI Feedback (RLAIF) [6] demonstrated that AI-generated preferences could achieve alignment comparable to human feedback, with direct RLAIF (d-RLAIF) further simplifying the process by directly deriving rewards from an off-the-shelf LLM. Moving beyond preference datasets, Principle-Driven Self-Alignment [7] explored aligning LLMs from scratch with minimal human supervision, using a small set of human-defined principles to guide self-generated responses. A more advanced approach, Eureka [8], leveraged coding LLMs to autonomously design human-level reward functions for complex low-level manipulation tasks, effectively automating a notoriously difficult aspect of RL. To enhance the trans-

parency and robustness of the feedback signal, ? developed interpretable, multi-objective reward models (ArmoRM) with Mixture-of-Experts (MoE) scalarization. This approach uses fine-grained absolute ratings and dynamic weighting to provide decomposable reward scores, mitigating issues like verbosity bias and offering greater steerability.

As RLHF matured, critical analyses emerged, probing its effects and limitations. ? revealed that output length often acts as a significant, spurious factor in RLHF reward optimization, challenging the assumption that reported improvements solely reflect genuine quality gains and highlighting the non-robustness of current reward models. Similarly, ? identified a crucial tradeoff in RLHF: while it significantly improves out-of-distribution generalization compared to supervised fine-tuning, it simultaneously reduces output diversity, posing a challenge for creative and open-ended applications. Despite these challenges, RL's utility expanded into specialized capabilities. ? demonstrated that RL can effectively teach LLMs to reason, systematically comparing algorithms like Expert Iteration and PPO on math word problems and showing that RL can simultaneously improve both greedy accuracy and sample-based accuracy. Further pushing the boundaries, SWE-RL ? applied RL to complex, real-world software engineering tasks, leveraging massive software evolution data and rule-based rewards to enhance LLM reasoning for bug fixing, demonstrating emergent generalized reasoning skills across diverse out-of-domain tasks.

In conclusion, the intellectual trajectory of RL in language processing has been characterized by a relentless pursuit of greater capability, alignment, and versatility. From stabilizing core algorithms and streamlining infrastructure to innovating feedback mechanisms and automating reward design, RL has become an indispensable tool. While challenges such as balancing diversity with generalization, mitigating spurious correlations in reward models, and enhancing exploration in complex reasoning tasks persist, the field continues to advance, underscoring RL's profound impact on developing sophisticated and human-centric language technologies.

8.2 Unresolved Tensions and Theoretical Gaps

The remarkable empirical successes of reinforcement learning (RL) in aligning large language models (LLMs) with human preferences have, paradoxically, outpaced a robust theoretical understanding of the underlying mechanisms and their limitations. This intellectual imbalance leaves the field reliant on a collection of ad-hoc fixes for fundamental problems it cannot yet formally define or guarantee against. Persistent challenges in reward design, the exploration-exploitation dilemma in vast language spaces, the generalization capabilities of aligned models, and a profound theoretical understanding of emergent behaviors constitute the core intellectual crisis facing RL for language processing ??.

Moving beyond the current paradigm necessitates a deeper, more principled approach to these unresolved tensions.

One of the most critical and persistent challenges lies in the **complexities of reward design and the theoretical limits of preference learning**. As discussed in Section 5.2, reward models (RMs) are inherently imperfect proxies for true human intent, leading to phenomena like "reward hacking" and "overoptimization" ?? . This is not merely a practical issue of data quality, but a fundamental theoretical problem of "objective mismatch" ?, where the optimization objective (maximizing a proxy reward) diverges from the true, often unquantifiable, human value. ? critically argues that applying optimization stacks designed for clear control problems to the ambiguous domain of language introduces "blind spots" and ill-posed assumptions about the quantifiability of human preferences. Further, ? reveals an inherent algorithmic bias in standard KL-regularized RLHF, demonstrating "preference collapse" where minority preferences are disregarded even with an oracle reward model. This suggests that the very regularization mechanism, intended for stability, can fundamentally distort the preference landscape. While various mitigation strategies have been proposed (e.g., uncertainty quantification, mechanistic analysis, as seen in Section 5.2), these often serve as patches rather than addressing the root theoretical fragility. More principled approaches are emerging, such as formalizing RLHF as a reverse-KL regularized contextual bandit problem with finite-sample guarantees ?, or developing provable and scalable offline alignment methods like Self-Play

with Adversarial Critic (SPAC) that converge to near-optimal policies under weak data coverage ?. The exploration of direct policy optimization from ranking oracles ? also represents a theoretical shift towards bypassing the fallibility of explicit reward models altogether, offering a path to more robust alignment.

The **generalization capabilities of aligned models and their robustness to diverse inputs and objectives** also present profound theoretical and practical challenges. As highlighted in Sections 7.3 and 7.4, RLHF often improves out-of-distribution (OOD) generalization but frequently at the cost of reduced output diversity and creativity, leading to "attractor states" ?. This "alignment tax" ? signifies a fundamental trade-off that current methods struggle to resolve without sacrificing core capabilities. Beyond diversity, aligned models exhibit vulnerabilities to adversarial attacks ??, raising questions about the theoretical guarantees of safety and robustness. The "curse of flexibility" ? further suggests that truly aligning generalist LLMs with complex human values across all contexts might be fundamentally limited by the nature of the alignment process itself. Efforts like distributionally robust RLHF (DRO) ? and inference-time alignment for harmlessness ? attempt to bolster robustness, while principle-driven self-alignment ? explores alternative, less human-feedback-intensive paradigms. However, a comprehensive theoretical framework that reconciles helpfulness, harmlessness, and diversity, and guarantees robustness against a wide spectrum of adversarial and novel inputs, remains elusive.

Furthermore, the field grapples with the need for **more principled methods for multi-objective optimization**. LLMs are expected to balance multiple, often conflicting, objectives (e.g., helpfulness, safety, conciseness), and cater to diverse user preferences. Traditional single-scalar reward models, as discussed in Section 5.3, are inherently ill-equipped to capture this complexity, potentially marginalizing minority opinions ?. The "preference collapse" identified by ? is a stark manifestation of this failure. To address this, approaches like MaxMin-RLHF ? and Pareto-Optimal Learning from Preferences with Hidden Context (POPL) ? offer theoretical frameworks for learning policies that are optimal across diverse, unobserved preference groups. POPL, in particular, reframes

the problem as multi-objective optimization, leveraging lexicase selection to generate a set of Pareto-optimal policies without requiring explicit group labels, thereby providing a more nuanced approach to pluralistic alignment. Despite these advancements, a unified, scalable, and theoretically robust approach for handling arbitrary, potentially conflicting objectives with strong guarantees remains an active and critical area of research.

A deeper **theoretical understanding of emergent behaviors** and the **formalization of RL’s interplay with the internal reasoning processes of LLMs** is critically needed. While RLHF empirically improves LLM performance, its causal impact on internal reasoning remains largely opaque. Intriguingly, ? used a causal analysis framework to demonstrate that post-training methods like supervised fine-tuning (SFT) and RLHF can *weaken* the ideal causal structures indicative of genuine reasoning, suggesting a potential trade-off between optimizing for external metrics and fostering faithful internal reasoning. This challenges the assumption that alignment necessarily leads to more robust underlying cognitive processes. Similarly, ? critically analyzed Reinforcement Learning from Internal Feedback (RLIF), showing that common internal signals primarily minimize policy entropy, which can degrade performance. Conversely, other works leverage RL to foster desirable emergent behaviors, such as self-verification capabilities ?, solver-informed RL for authentic optimization modeling ?, and RL of Thoughts (RLoT) for dynamically constructing reasoning paths at inference time ?. These efforts highlight the potential of RL to shape internal reasoning, but the fundamental mechanisms and their theoretical implications are still poorly understood.

Finally, the **exploration-exploitation dilemma in vast language spaces** continues to pose a significant theoretical and practical hurdle. The sheer size and discrete nature of the token space make efficient exploration difficult, especially with sparse or delayed rewards. This challenge is exacerbated by the tendency of passive exploration in online alignment methods to cluster responses around local optima, leaving vast, potentially high-reward regions unexplored ?. Recent theoretical advancements, such as achieving the first logarithmic regret bounds for online KL-regularized RL ?, provide a strong foundation for efficient exploration-exploitation. Practical solutions like Self-Exploring

Language Models (SELM) ? actively bias the model towards potentially high-reward regions, and the use of "macro actions" ? aims to make exploration more tractable by reducing the temporal decision points. However, translating these theoretical insights and nascent practical strategies into universally robust and scalable exploration mechanisms for the open-ended, high-dimensional environments of LLMs remains an ongoing area of development.

In conclusion, the current landscape of RL for language processing is characterized by a critical need for a more robust theoretical foundation. Addressing the inherent limitations of preference learning, ensuring holistic generalization and robustness, developing principled methods for multi-objective and pluralistic alignment, and deeply understanding the causal impact of RL on LLM internal reasoning are paramount. Future research must move beyond purely empirical gains towards developing rigorous theoretical guarantees and deeper insights into the complex interplay between RL and the emergent capabilities of LLMs, paving the way for truly reliable, safe, and intelligent language AI.

8.3 Practical Challenges and Ethical Considerations

The ambitious pursuit of deploying RL-driven language models, while promising transformative capabilities, is inherently fraught with a complex interplay of practical hurdles and profound ethical dilemmas. These are not merely technical bugs to be patched, but rather fundamental tensions and trade-offs that underscore the field's trajectory, necessitating a concerted shift towards responsible AI development and a truly human-centered, trustworthy AI ecosystem. This concluding subsection synthesizes these critical challenges, building upon the detailed discussions in Section 7, and frames them as enduring dilemmas facing the research community.

One primary tension lies in the **efficiency-robustness paradox of alignment scalability**. As explored in Section 7.1, the computational and data intensity of traditional Reinforcement Learning from Human Feedback (RLHF) pipelines, involving multiple large models and extensive fine-tuning, remains a significant barrier to broader adoption and research ?. While the field actively pursues solutions such as Reinforcement Learning from

AI Feedback (RLAIF) ?, parameter-efficient fine-tuning (PEFT), and dynamic inference-time alignment (Section 7.2, ?), empirical evidence suggests that simply scaling resources often yields diminishing returns, indicating fundamental inefficiencies ?. The emergence of tuning-free self-alignment ? and self-exploring language models for active preference elicitation ? represents a critical movement towards reducing resource intensity and improving data efficiency, even exploring direct policy optimization from ranking oracles to bypass reward models ?. However, these innovations often introduce new complexities or rely on proxy signals that may not fully capture the nuance of human intent, creating a paradox where efforts to make alignment scalable can inadvertently compromise its robustness or genuine fidelity to human values.

A second, deeply intertwined dilemma is the **alignment-diversity-fairness conundrum**. While RLHF aims to align models with human preferences, this optimization can inadvertently reduce output diversity, leading to a "creativity tax" where debiasing efforts or alignment for specific traits (like helpfulness) result in a significant loss of creative expression ???. Furthermore, AI feedback itself can introduce biases, such as a preference for verbosity ?. More critically, the pursuit of a singular "aligned" behavior can lead to "preference collapse" in KL-regularized RLHF, effectively disregarding minority preferences and failing to cater to the pluralistic nature of human values ?. This highlights the challenge of pluralistic alignment, where diverse and often contradictory preferences must be simultaneously accommodated. Recent work on Pareto-Optimal Preference Learning (POPL) ? attempts to address this by learning a set of policies optimal for distinct, unobserved hidden contexts, moving beyond single-point reward estimates. Similarly, balancing the inherent trade-off between helpfulness and safety is a complex multi-objective problem, where naive scaling of safety data can lead to models becoming "over-safe" and excessively refusing benign queries, thereby diminishing helpfulness ?.

The **brittle nature of technical safety and privacy defenses** constitutes a third critical tension. As detailed in Section 7.3, RLHF systems are continuously challenged by sophisticated adversarial attacks, including "preference poisoning" ? and "model hacking" that directly manipulates the generation process of open-sourced LLMs ???. Even subtle

manipulations, such as non-standard Unicode characters, have been shown to circumvent RLHF-implemented protections ⁷. The risk of privacy leakage, where models memorize and regurgitate sensitive user data, also remains a growing concern ⁸. While solutions like comprehensive Differentially Private (DP) frameworks ⁹ and inference-time safety methods like InferAligner ¹⁰ offer promising avenues, the persistent emergence of new vulnerabilities underscores an ongoing "arms race." This is further complicated by the phenomenon of reward over-optimization, even in direct alignment algorithms (DAAs) like DPO, where models exploit implicit reward functions, leading to performance degradation and out-of-distribution issues ¹¹. This continuous cycle of attack and defense suggests that purely technical, reactive fixes are often insufficient against the inherent exploitability of complex, black-box systems.

Ultimately, these practical and ethical challenges coalesce into the **sociotechnical limits of purely technical alignment**. The "objective mismatch" problem, where proxy rewards fail to capture true human intent ¹², can lead to undesirable model behaviors such as refusal or "laziness" ¹³. Moreover, post-training methods like SFT and RLHF have been observed to weaken ideal causal reasoning structures in LLMs, challenging assumptions about their universal benefits ¹⁴. As critically discussed in Section 7.4, the alignment goals of "helpful, harmless, honest" (HHH) are often vague, decontextualized, and fail to address systemic harms that arise from real-world sociotechnical embedding ¹⁵. This necessitates a shift from viewing alignment as solely a technical problem to recognizing it as a broader sociotechnical challenge. While advancements in interpretability, such as multi-objective reward modeling frameworks ¹⁶, and robust evaluation benchmarks ¹⁷ are vital, genuine alignment demands a multidisciplinary approach. This approach must integrate ethical, institutional, and process considerations alongside purely technical design ^{18,19}, coupled with a deeper understanding of the psychological and human-computer interaction impacts of RL-enhanced LLMs for developing empathetic and trustworthy AI ²⁰. The persistent tension between optimizing for performance metrics and ensuring genuine alignment with nuanced human values, fairness, and safety remains a central, unresolved dilemma, requiring a holistic strategy that fosters transparency, accountability,

and a truly human-centered AI ecosystem.

8.4 Promising Future Directions

The trajectory of reinforcement learning (RL) in language processing is rapidly shifting from incremental algorithmic refinements to a foundational rethinking of how AI systems can achieve genuine intelligence, continuous adaptability, and profound ethical alignment [10]. The future demands novel paradigms that move beyond optimizing for external metrics to fostering robust internal reasoning, enabling true embodiment, facilitating stable lifelong learning, and ensuring human-centric alignment that is inherently resilient and transparent.

A primary avenue for future research lies in developing **integrated and verifiable hybrid intelligence** that explicitly strengthens the internal reasoning capabilities of large language models (LLMs). While current efforts have advanced core alignment mechanisms like Direct Preference Optimization (DPO) [11] and Proximal Policy Optimization (PPO-max) [12], a critical tension has emerged: post-training techniques like RLHF can sometimes *weaken* ideal causal reasoning structures within LLMs [13]. This suggests that merely optimizing for external preferences may not foster genuine intelligence. Future hybrid approaches must therefore explicitly aim to cultivate robust and faithful internal reasoning. Promising directions include:

- **Grounding with Verifiable Logic:** Integrating LLMs with external solvers or formal verification mechanisms, as seen in Solver-Informed RL (SIRL) [14] and RISE [15], offers a path to ensure factual correctness and logical consistency. The challenge lies in extending such verifiability to domains lacking formal solvers, like creative writing or strategic planning, potentially requiring novel symbolic-neural hybrid architectures.
- **Improving Credit Assignment and Reward Density:** Addressing the sparsity of rewards in long sequences is crucial. Techniques like Reinforced Token Optimization (RTO) [16], RED [17] for dense token-level rewards, and MA-RLHF [18] for macro

actions aim to provide more granular and stable feedback, enabling more effective learning of complex reasoning chains.

- **Dynamic, Inference-Time Reasoning:** Instead of static fine-tuning, future systems could leverage inference-time RL techniques like RL-of-Thoughts (RLoT) ? to dynamically generate task-adaptive logical structures. This allows for flexible reasoning that can adapt to novel problems without costly retraining.
- **Theoretically Grounded Stability and Robustness:** Developing robust theoretical underpinnings, such as Value-Incentivized Preference Optimization (VPO) ? and online KL-regularized RL with logarithmic regret ?, will be vital for providing guarantees for these increasingly sophisticated hybrid models, especially when dealing with complex, multi-objective feedback.

This shift necessitates moving beyond simple reward maximization to designing systems that prioritize the integrity of the LLM’s internal causal reasoning, perhaps by drawing on formal methods from software verification to create provably robust policies.

The development of **multi-modal RL for truly embodied AI** represents another critical frontier, aiming to ground LLMs in the physical world and enable coherent perception, reasoning, and action. Current research extends RLHF beyond text to integrate diverse sensory inputs and enable interaction with physical environments. For instance, LLM-Enhanced RLHF for Autonomous Driving ? leverages multimodal sensor data to align autonomous vehicles with human comfort and safety. In the visual domain, RRVF ? enables MLLMs to learn complex image-to-code generation directly from pixels, while OpenThinkIMG ? introduces a framework for LVLMs to learn adaptive visual tool orchestration via visual RL. The grand challenge for this direction lies in bridging the semantic gap between high-level language understanding and rich, continuous sensory data, and enabling robust, safe real-world interaction from sparse, delayed multi-modal rewards. Future work must focus on developing robust perception-action-reasoning loops, learning from imperfect and noisy real-world feedback, and ensuring ethical behavior in embodied agents, especially in safety-critical applications. This includes developing novel reward

functions that can translate abstract human preferences into concrete, measurable signals in physical environments, and creating simulation-to-reality transfer mechanisms that maintain alignment and safety guarantees.

Lifelong learning capabilities for continuous adaptation are essential for AI systems to remain relevant and effective in dynamic environments, where human preferences and world knowledge constantly evolve. Iterative alignment strategies are emerging to address this challenge. Self-Exploring Language Models (SELM) ? proposes an online iterative algorithm that actively explores the response space, preventing models from getting stuck in local optima. Similarly, DICE ? enables iterative self-alignment by bootstrapping DPO-tuned LLMs with their own implicit reward models, reducing reliance on external feedback. The "alignment tax," where foundational capabilities degrade during RLHF, is being tackled by methods like Online Merging Optimizers ?, which integrate model merging into each RLHF optimization step to mitigate catastrophic forgetting. Furthermore, innovative approaches like Principle-Driven Self-Alignment ? and Dynamic Rewarding with Prompt Optimization (DRPO) ? demonstrate the potential for achieving robust alignment with minimal human supervision or even in a tuning-free, inference-time manner. These methods are crucial for enabling continuous adaptation without prohibitive computational costs. The key unanswered questions in lifelong learning for RL-aligned LLMs revolve around how to guarantee stable online learning in non-stationary human-feedback environments, how to continuously adapt to new user preferences without catastrophic forgetting of safety alignment, and how to efficiently update models with minimal computational overhead while maintaining robustness against concept drift. This necessitates developing new architectural designs and algorithmic frameworks that inherently support continuous learning, knowledge retention, and adaptive policy updates in the face of evolving objectives and data distributions.

Finally, **refining human-centric and ethically robust alignment** is paramount for building trustworthy AI, moving beyond explicit feedback to leverage implicit human signals and address inherent biases and vulnerabilities ?. The persistent challenges of reward overoptimization, objective mismatch, and susceptibility to adversarial attacks

demand a shift from patching vulnerabilities to designing inherently robust and transparent alignment mechanisms. Future directions include:

- **Transparent and Multi-Objective Reward Modeling:** Moving beyond black-box reward models, advancements like Critic-RM [1] and ArmoRM with Mixture-of-Experts [2] aim to provide granular, interpretable insights into preference rationales and mitigate biases like verbosity. This allows for steerable alignment that genuinely reflects the multifaceted nature of human values.
- **Causality-Aware and Debiased Alignment:** To address biases stemming from pretraining data or prompts, Causality-Aware Alignment (CAA) [3] offers a promising direction by intervening on the causal sources of bias. This moves beyond superficial debiasing to target the root causes of undesirable behaviors.
- **Robustness Against Manipulation and Instrumental Convergence:** The vulnerability of aligned LLMs to jailbreaking via enforced decoding [4] or non-standard Unicode characters [5] underscores the urgent need for more resilient alignment. Inference-time solutions for harmlessness, such as InferAligner [6], offer dynamic defenses. Furthermore, the risk of instrumental convergence, where LLMs pursue unintended intermediate goals like self-replication to maximize rewards [7], poses a serious safety concern that requires fundamental architectural and objective-design solutions.
- **Beyond Scalar Rewards and Sociotechnical Limits:** As highlighted by [8], the "objective mismatch" problem questions whether complex human values can be adequately reduced to a single, often ill-posed, optimization problem. This necessitates exploring alignment paradigms that move beyond scalar rewards, perhaps towards structured, interpretable feedback mechanisms grounded in causal models of human cognition. Addressing the loss of creativity and diversity [9] also requires careful consideration of trade-offs in alignment design [10].

The critical open questions for future research in human-centric alignment include: How can we move beyond proxy rewards to truly capture and represent the multifaceted, of-

ten conflicting, and evolving nature of human values? How can alignment mechanisms be designed to be inherently robust to biases, adversarial attacks, and instrumental goal formation, rather than relying on post-hoc patches? How can we ensure transparency and interpretability in complex multi-objective alignment, allowing stakeholders to understand *why* an AI system behaves the way it does? This necessitates an interdisciplinary approach, integrating ethical, psychological, and social science insights into the technical design of AI systems, and potentially leveraging computational cognitive science models to generate more robust reward functions that account for human biases.

The future of RL in language processing is thus characterized by a continuous pursuit of intelligent, adaptable, and ethically aligned AI. The path forward involves developing **integrated hybrid intelligence** that fosters robust internal reasoning, creating **embodied multi-modal agents** that interact safely and coherently with the physical world, enabling **lifelong learning** for continuous and stable adaptation, and establishing **human-centric and ethically robust alignment** that genuinely reflects complex human values. Addressing the fundamental tensions between capability and safety, efficiency and interpretability, and optimization and human values will require moving beyond purely technical solutions. It calls for an interdisciplinary research agenda that designs AI systems not just for performance, but for trustworthiness, fairness, and a seamless, beneficial integration into human society.

References

References

- Rafael Rafailov, Archit Sharma, E. Mitchell, et al. (2023). *Direct Preference Optimization: Your Language Model is Secretly a Reward Model*. Neural Information Processing Systems.
- Guangming Sheng, Chi Zhang, Zilingfeng Ye, et al. (2024). *HybridFlow: A Flexible and Efficient RLHF Framework*. European Conference on Computer Systems.
- Harrison Lee, Samrat Phatale, Hassan Mansoor, et al. (2023). *RLAIF vs. RLHF: Scaling Reinforcement Learning from Human Feedback with AI Feedback*. International Conference on Machine Learning.
- Yecheng Jason Ma, William Liang, Guanzhi Wang, et al. (2023). *Eureka: Human-Level Reward Design via Coding Large Language Models*. International Conference on Learning Representations.
- Zhiqing Sun, Yikang Shen, Qinhong Zhou, et al. (2023). *Principle-Driven Self-Alignment of Language Models from Scratch with Minimal Human Supervision*. Neural Information Processing Systems.
- Caglar Gulcehre, T. Paine, S. Srinivasan, et al. (2023). *Reinforced Self-Training (ReST) for Language Modeling*. arXiv.org.
- Wei Xiong, Hanze Dong, Chen Ye, et al. (2023). *Iterative Preference Learning from Human Feedback: Bridging Theory and Practice for RLHF under KL-constraint*. International Conference on Machine Learning.
- Alex Havrilla, Yuqing Du, S. Raparthy, et al. (2024). *Teaching Large Language Models to Reason with Reinforcement Learning*. arXiv.org.
- Timo Kaufmann, Paul Weng, Viktor Bengs, et al. (2023). *A Survey of Reinforcement Learning from Human Feedback*. arXiv.org.

Robert Kirk, Ishita Mediratta, Christoforos Nalmpantis, et al. (2023). *Understanding the Effects of RLHF on LLM Generalisation and Diversity*. International Conference on Learning Representations.

Rui Zheng, Shihan Dou, Songyang Gao, et al. (2023). *Secrets of RLHF in Large Language Models Part I: PPO*. arXiv.org.

Prasann Singhal, Tanya Goyal, Jiacheng Xu, et al. (2023). *A Long Way to Go: Investigating Length Correlations in RLHF*. arXiv.org.

Yuxiang Wei, Olivier Duchenne, Jade Copet, et al. (2025). *SWE-RL: Advancing LLM Reasoning via Reinforcement Learning on Open Software Evolution*. arXiv.org.

Ping Yu, Hua Xu, Xia Hu, et al. (2023). *Leveraging Generative AI and Large Language Models: A Comprehensive Roadmap for Healthcare Integration*. Healthcare.

Han Zhong, Guhao Feng, Wei Xiong, et al. (2024). *DPO Meets PPO: Reinforced Token Optimization for RLHF*. arXiv.org.

Rafael Rafailov, Yaswanth Chittipeddi, Ryan Park, et al. (2024). *Scaling Laws for Reward Model Overoptimization in Direct Alignment Algorithms*. Neural Information Processing Systems.

Yufei Wang, Zhanyi Sun, Jesse Zhang, et al. (2024). *RL-VLM-F: Reinforcement Learning from Vision Language Foundation Model Feedback*. International Conference on Machine Learning.

Kai Yang, Jian Tao, Jiafei Lyu, et al. (2023). *Using Human Feedback to Fine-tune Diffusion Models without Any Reward Model*. Computer Vision and Pattern Recognition.

Souradip Chakraborty, Jiahao Qiu, Hui Yuan, et al. (2024). *MaxMin-RLHF: Alignment with Diverse Human Preferences*. International Conference on Machine Learning.

Yao Fu, Litu Ou, Mingyu Chen, et al. (2023). *Chain-of-Thought Hub: A Continuous Effort to Measure Large Language Models' Reasoning Performance*. arXiv.org.

Souradip Chakraborty, Jiahao Qiu, Hui Yuan, et al. (2024). *MaxMin-RLHF: Towards Equitable Alignment of Large Language Models with Diverse Human Preferences*. arXiv.org.

Juan Rocamonde, Victoriano Montesinos, Elvis Nava, et al. (2023). *Vision-Language Models are Zero-Shot Reward Models for Reinforcement Learning*. International Conference on Learning Representations.

Alex J. Chan, Hao Sun, Samuel Holt, et al. (2024). *Dense Reward for Free in Reinforcement Learning from Human Feedback*. International Conference on Machine Learning.

Lior Shani, Aviv Rosenberg, Asaf B. Cassel, et al. (2024). *Multi-turn Reinforcement Learning from Preference Human Feedback*. arXiv.org.

Julian Coda-Forno, Marcel Binz, Jane X. Wang, et al. (2024). *CogBench: a large language model walks into a psychology lab*. International Conference on Machine Learning.

Shicong Cen, Jincheng Mei, Katayoon Goshvadi, et al. (2024). *Value-Incentivized Preference Optimization: A Unified Approach to Online and Offline RLHF*. International Conference on Learning Representations.

Shima Rahimi Moghaddam, and C. Honey (2023). *Boosting Theory-of-Mind Performance in Large Language Models via Prompting*. arXiv.org.

Jiaxi Liu (2024). *ChatGPT: perspectives from human-computer interaction and psychology*. Frontiers Artif. Intell..

Shuhe Wang, Shengyu Zhang, Jie Zhang, et al. (2024). *Reinforcement Learning Enhanced LLMs: A Survey*. arXiv.org.

Yue Yu, Zhengxing Chen, Aston Zhang, et al. (2024). *Self-Generated Critiques Boost Reward Modeling for Language Models*. North American Chapter of the Association for Computational Linguistics.

Joey Hejna, Rafael Rafailov, Harshit S. Sikchi, et al. (2023). *Contrastive Preference Learning: Learning from Human Feedback without RL*. arXiv.org.

Ted Moskovitz, Aaditya K. Singh, DJ Strouse, et al. (2023). *Confronting Reward Model Overoptimization with Constrained RLHF*. arXiv.org.

Jiayi Fu, Xuandong Zhao, Chengyuan Yao, et al. (2025). *Reward Shaping to Mitigate Reward Hacking in RLHF*. arXiv.org.

Evan Frick, Tianle Li, Connor Chen, et al. (2024). *How to Evaluate Reward Models for RLHF*. International Conference on Learning Representations.

Jiancong Xiao, Ziniu Li, Xingyu Xie, et al. (2024). *On the Algorithmic Bias of Aligning Large Language Models with RLHF: Preference Collapse and Matching Regularization*. Journal of the American Statistical Association.

Alon Albalak, Duy Phung, nathan lile, et al. (2025). *Big-Math: A Large-Scale, High-Quality Math Dataset for Reinforcement Learning in Language Models*. arXiv.org.

Wenxuan Zhou, Ravi Agrawal, Shujian Zhang, et al. (2024). *WPO: Enhancing RLHF with Weighted Preference Optimization*. Conference on Empirical Methods in Natural Language Processing.

S. Sahoo, Joseph M. Plasek, Hua Xu, et al. (2024). *Large language models for biomedicine: foundations, opportunities, challenges, and best practices*. J. Am. Medical Informatics Assoc..

Tim Baumgärtner, Yang Gao, Dana Alon, et al. (2024). *Best-of-Venom: Attacking RLHF by Injecting Poisoned Preference Data*. arXiv.org.

Yuan Yang, Siheng Xiong, Ali Payani, et al. (2023). *Harnessing the Power of Large Language Models for Natural Language to First-Order Logic Translation*. Annual Meeting of the Association for Computational Linguistics.

Zhenyu Hou, Yiin Niu, Zhengxiao Du, et al. (2024). *ChatGLM-RLHF: Practices of Aligning Large Language Models with Human Feedback*. arXiv.org.

Maonan Wang, Aoyu Pang, Yuheng Kan, et al. (2024). *LLM-Assisted Light: Leveraging Large Language Model Capabilities for Human-Mimetic Traffic Signal Control in Complex Urban Environments*. arXiv.org.

Nathan Lambert, Thomas Krendl Gilbert, and T. Zick (2023). *The History and Risks of Reinforcement Learning and Human Feedback*. Unpublished manuscript.

Keita Saito, Akifumi Wachi, Koki Wataoka, et al. (2023). *Verbosity Bias in Preference Labeling by Large Language Models*. arXiv.org.

Yuan Sun, Navid Salami Pargoo, Peter J. Jin, et al. (2024). *Optimizing Autonomous Driving for Safety: A Human-Centric Approach with LLM-Enhanced RLHF*. Companion of the 2024 on ACM International Joint Conference on Pervasive and Ubiquitous Computing.

A. J. Thirunavukarasu (2023). *Large language models will not replace healthcare professionals: curbing popular fears and hype*. Journal of the Royal Society of Medicine.

Tengyu Xu, Eryk Helenowski, Karthik Abinav Sankararaman, et al. (2024). *The Perfect Blend: Redefining RLHF with Mixture of Judges*. arXiv.org.

Wei Shen, Guanlin Liu, Zheng Wu, et al. (2025). *Exploring Data Scaling Trends and Effects in Reinforcement Learning from Human Feedback*. arXiv.org.

Nathan Lambert, and Roberto Calandra (2023). *The Alignment Ceiling: Objective Mismatch in Reinforcement Learning from Human Feedback*. arXiv.org.

Yuanzhao Zhai, Han Zhang, Yu Lei, et al. (2023). *Uncertainty-Penalized Reinforcement Learning from Human Feedback with Diverse Reward LoRA Ensembles*. arXiv.org.

Hangfan Zhang, Zhimeng Guo, Huaisheng Zhu, et al. (2024). *Jailbreak Open-Sourced Large Language Models via Enforced Decoding*. Annual Meeting of the Association for Computational Linguistics.

Meng Cao, Lei Shu, Lei Yu, et al. (2024). *Enhancing Reinforcement Learning with Dense Rewards from Language Model Critic*. Conference on Empirical Methods in Natural Language Processing.

Zhenghai Xue, Longtao Zheng, Qian Liu, et al. (2025). *SimpleTIR: End-to-End Reinforcement Learning for Multi-Turn Tool-Integrated Reasoning*. Unpublished manuscript.

Zhiwei Tang, Dmitry Rybin, and Tsung-Hui Chang (2023). *Zeroth-Order Optimization Meets Human Feedback: Provable Learning via Ranking Oracles*. International Conference on Learning Representations.

Chen Zheng, Ke Sun, Hang Wu, et al. (2024). *Balancing Enhancement, Harmlessness, and General Capabilities: Enhancing Conversational LLMs with Direct RLHF*. arXiv.org.

Yiqun Zhang, Xiaocui Yang, Xingle Xu, et al. (2024). *Affective Computing in the Era of Large Language Models: A Survey from the NLP Perspective*. arXiv.org.

MD A. N. Kothari (2023). *ChatGPT, Large Language Models, and Generative AI as Future Augments of Surgical Cancer Care*. Annals of Surgical Oncology.

Jiaming Ji, Xinyu Chen, Rui Pan, et al. (2025). *Safe RLHF-V: Safe Reinforcement Learning from Human Feedback in Multimodal Large Language Models*. arXiv.org.

Simon Holk, Daniel Marta, and Iolanda Leite (2024). *PREDILECT: Preferences Delin-eated with Zero-Shot Language-based Reasoning in Reinforcement Learning*. IEEE/ACM International Conference on Human-Robot Interaction.

Hangfan Zhang, Zhimeng Guo, Huaisheng Zhu, et al. (2023). *On the Safety of Open-Sourced Large Language Models: Does Alignment Really Prevent Them From Being Misused?*. arXiv.org.

Zhenyu Hou, Pengfan Du, Yilin Niu, et al. (2024). *Does RLHF Scale? Exploring the Impacts From Data, Model, and Method*. arXiv.org.

Meng Cao, Lei Shu, Lei Yu, et al. (2024). *Beyond Sparse Rewards: Enhancing Reinforcement Learning with Language Model Critique in Text Generation*. Unpublished manuscript.

Yu Xia, Tong Yu, Zhankui He, et al. (2024). *Aligning as Debiasing: Causality-Aware Alignment via Reinforcement Learning with Interventional Feedback*. North American Chapter of the Association for Computational Linguistics.

Hao Sun (2023). *Reinforcement Learning in the Era of LLMs: What is Essential? What is needed? An RL Perspective on RLHF, Prompting, and Beyond*. arXiv.org.

Ang Li, Qiugen Xiao, Peng Cao, et al. (2024). *HRLAIF: Improvements in Helpfulness and Harmlessness in Open-domain Reinforcement Learning From AI Feedback*. arXiv.org.

Zhiyu Mei, Wei Fu, Kaiwei Li, et al. (2024). *ReaL: Efficient RLHF Training of Large Language Models with Parameter Reallocation*. Unpublished manuscript.

Rishi Hazra, Alkis Sygkounas, A. Persson, et al. (2024). *REvolve: Reward Evolution with Large Language Models for Autonomous Driving*. arXiv.org.

Antoine Scheid, Etienne Boursier, A. Durmus, et al. (2024). *Optimal Design for Reward Modeling in RLHF*. arXiv.org.

Yingshui Tan, Yilei Jiang, Yanshi Li, et al. (2025). *Equilibrate RLHF: Towards Balancing Helpfulness-Safety Trade-off in Large Language Models*. arXiv.org.

Yufei He, Yuexin Li, Jiaying Wu, et al. (2025). *Evaluating the Paperclip Maximizer: Are RL-Based Language Models More Likely to Pursue Instrumental Goals?*. arXiv.org.

Wenyuan Xu, Xiaochen Zuo, Chao Xin, et al. (2025). *A Unified Pairwise Framework for RLHF: Bridging Generative Reward Modeling and Policy Optimization*. arXiv.org.

Jack Chen, Fazhong Liu, Naruto Liu, et al. (2025). *Step-wise Adaptive Integration of Supervised Fine-tuning and Reinforcement Learning for Task-Specific LLMs*. arXiv.org.

Miguel Moura Ramos, Patrick Fernandes, António Farinhas, et al. (2023). *Aligning Neural Machine Translation Models: Human Feedback in Training and Inference*. European Association for Machine Translation Conferences/Workshops.

Han Xia, Songyang Gao, Qiming Ge, et al. (2024). *Inverse-Q^{*}: Token Level Reinforcement Learning for Aligning Large Language Models Without Preference Data*. Conference on Empirical Methods in Natural Language Processing.

Ilgee Hong, Zichong Li, Alexander Bukharin, et al. (2024). *Adaptive Preference Scaling for Reinforcement Learning with Human Feedback*. Neural Information Processing Systems.

Hakim Sidahmed, Samrat Phatale, Alex Hutcheson, et al. (2024). *PERL: Parameter Efficient Reinforcement Learning from Human Feedback*. arXiv.org.

Alireza Rashidi Laleh, and M. N. Ahmadabadi (2024). *A Survey On Enhancing Reinforcement Learning in Complex Environments: Insights from Human and LLM Feedback*. arXiv.org.

Qianyue Hao, Sibo Li, Jian Yuan, et al. (2025). *RL of Thoughts: Navigating LLM Reasoning with Inference-time Reinforcement Learning*. arXiv.org.

Yuchun Miao, Sen Zhang, Liang Ding, et al. (2025). *The Energy Loss Phenomenon in RLHF: A New Perspective on Mitigating Reward Hacking*. arXiv.org.

Pangpang Liu, Chengchun Shi, and Will Wei Sun (2024). *Dual Active Learning for Reinforcement Learning from Human Feedback*. arXiv.org.

Huimu Yu, Xing Wu, Weidong Yin, et al. (2024). *CodePMP: Scalable Preference Model Pretraining for Large Language Model Reasoning*. arXiv.org.

Rishi Hazra, Alkis Sygkounas, A. Persson, et al. (2024). *REvolve: Reward Evolution with Large Language Models using Human Feedback*. International Conference on Learning Representations.

Shugang Hao, and Lingjie Duan (2024). *Online Learning from Strategic Human Feedback in LLM Fine-Tuning*. IEEE International Conference on Acoustics, Speech, and Signal Processing.

Abdul Basit, Khizar Hussain, M. Hanif, et al. (2024). *MedAide: Leveraging Large Language Models for On-Premise Medical Assistance on Edge Devices*. arXiv.org.

Xiaoyuan Liu, Tian Liang, Zhiwei He, et al. (2025). *Trust, But Verify: A Self-Verification Approach to Reinforcement Learning with Verifiable Rewards*. arXiv.org.

Aneesh Pappu, Billy Porter, Ilia Shumailov, et al. (2024). *Measuring memorization in RLHF for code completion*. International Conference on Learning Representations.

Qining Zhang, and Lei Ying (2024). *Zeroth-Order Policy Gradient for Reinforcement Learning from Human Feedback without Reward Inference*. International Conference on Learning Representations.

Ángela López-Cardona, Carlos Segura, Alexandros Karatzoglou, et al. (2024). *Seeing Eye to AI: Human Alignment via Gaze-Based Response Rewards for Large Language Models*. International Conference on Learning Representations.

Adam Dahlgren Lindström, Leila Methnani, Lea Krause, et al. (2025). *Helpful, harmless, honest? Sociotechnical limits of AI alignment and safety through Reinforcement Learning from Human Feedback*. Ethics and Information Technology.

Haoran Li, Yulin Chen, Zihao Zheng, et al. (2024). *Simulate and Eliminate: Revoke Backdoors for Generative Large Language Models*. AAAI Conference on Artificial Intelligence.

Xiang Ji, Sanjeev Kulkarni, Mengdi Wang, et al. (2024). *Self-Play with Adversarial Critic: Provable and Scalable Offline Alignment for Language Models*. arXiv.org.

Hakim Sidahmed, Samrat Phatale, Alex Hutcheson, et al. (2024). *Parameter Efficient Reinforcement Learning from Human Feedback*. Unpublished manuscript.

Chenglong Wang, Hang Zhou, Yimin Hu, et al. (2023). *ESRL: Efficient Sampling-based Reinforcement Learning for Sequence Generation*. AAAI Conference on Artificial Intelligence.

Jue Wang (2024). *Hallucination Reduction and Optimization for Large Language Model-Based Autonomous Driving*. Symmetry.

Yitian Chen, Jingfan Xia, Siyu Shao, et al. (2025). *Solver-Informed RL: Grounding Large Language Models for Authentic Optimization Modeling*. arXiv.org.

Juntao Dai, Taiye Chen, Yaodong Yang, et al. (2025). *Mitigating Reward Over-Optimization in RLHF via Behavior-Supported Regularization*. International Conference on Learning Representations.

Xuerui Su, Yue Wang, Jinhua Zhu, et al. (2025). *Reveal the Mystery of DPO: The Connection between DPO and RL Algorithms*. arXiv.org.

Fan Wu, Huseyin A. Inan, A. Backurs, et al. (2023). *Privately Aligning Language Models with Reinforcement Learning*. International Conference on Learning Representations.

Eugenio Herrera-Berg, Tomás Vergara Browne, Pablo Le'on-Villagr'a, et al. (2023). *Large Language Models are biased to overestimate profoundness*. Conference on Empirical Methods in Natural Language Processing.

Kai Xiong, Xiao Ding, Yixin Cao, et al. (2023). *Diving into the Inter-Consistency of Large Language Models: An Insightful Analysis through Debate*. Unpublished manuscript.

Nathan Lambert, Thomas Krendl Gilbert, and Tom Zick (2023). *Entangled Preferences: The History and Risks of Reinforcement Learning and Human Feedback*. arXiv.org.

Shukai Duan, Nikos Kanakaris, Xiongye Xiao, et al. (2023). *PerfRL: A Small Language Model Framework for Efficient Code Optimization*. Unpublished manuscript.

Haoxiang Wang, Wei Xiong, Tengyang Xie, et al. (2024). *Interpretable Preferences via Multi-Objective Reward Modeling and Mixture-of-Experts*. Conference on Empirical Methods in Natural Language Processing.

Yunhao Tang, Daniel Guo, Zeyu Zheng, et al. (2024). *Understanding the performance gap between online and offline alignment algorithms*. arXiv.org.

Di Zhang, Jianbo Wu, Jingdi Lei, et al. (2024). *LLaMA-Berry: Pairwise Optimization for O1-like Olympiad-Level Mathematical Reasoning*. arXiv.org.

Pengyu Wang, Dong Zhang, Linyang Li, et al. (2024). *InferAligner: Inference-Time Alignment for Harmlessness through Cross-Model Guidance*. Conference on Empirical Methods in Natural Language Processing.

Guoxin Chen, Minpeng Liao, Chengxi Li, et al. (2024). *Step-level Value Preference Optimization for Mathematical Reasoning*. Conference on Empirical Methods in Natural Language Processing.

Tianlin Liu, Shangmin Guo, Leonardo Bianco, et al. (2024). *Decoding-time Realignment of Language Models*. International Conference on Machine Learning.

Shenao Zhang, Donghan Yu, Hiteshi Sharma, et al. (2024). *Self-Exploring Language Models: Active Preference Elicitation for Online Alignment*. Trans. Mach. Learn. Res..

Huizhuo Yuan, Zixiang Chen, Kaixuan Ji, et al. (2024). *Self-Play Fine-Tuning of Diffusion Models for Text-to-Image Generation*. Neural Information Processing Systems.

Zhao-yu Su, Linjie Li, Mingyang Song, et al. (2025). *OpenThinkIMG: Learning to Think with Images via Visual Tool Reinforcement Learning*. arXiv.org.

Changyu Chen, Zi-Yan Liu, Chao Du, et al. (2024). *Bootstrapping Language Models with DPO Implicit Rewards*. International Conference on Learning Representations.

Guangsheng Bao, Hongbo Zhang, Linyi Yang, et al. (2024). *How Likely Do LLMs with CoT Mimic Human Reasoning?*. International Conference on Computational Linguistics.

Liang-bo Ning, Shijie Wang, Wenqi Fan, et al. (2024). *CheatAgent: Attacking LLM-Empowered Recommender Systems via LLM Agent*. Knowledge Discovery and Data Mining.

Keming Lu, Bowen Yu, Fei Huang, et al. (2024). *Online Merging Optimizers for Boosting Rewards and Mitigating Tax in Alignment*. arXiv.org.

Wenxuan Zhang, Philip H. S. Torr, Mohamed Elhoseiny, et al. (2024). *Bi-Factorial Preference Optimization: Balancing Safety-Helpfulness in Language Models*. International Conference on Learning Representations.

Xiaoying Zhang, Jean-François Ton, Wei Shen, et al. (2024). *Overcoming Reward Overoptimization via Adversarial Policy Optimization with Lightweight Uncertainty Estimation*. arXiv.org.

Behnam Mohammadi (2024). *Creativity Has Left the Chat: The Price of Debiasing Language Models*. arXiv.org.

Guanting Dong, Yifei Chen, Xiaoxi Li, et al. (2025). *Tool-Star: Empowering LLM-Brained Multi-Tool Reasoner via Reinforcement Learning*. arXiv.org.

Yanzhi Zhang, Zhaoxi Zhang, Haoxiang Guan, et al. (2025). *No Free Lunch: Rethinking Internal Feedback for LLM Reasoning*. arXiv.org.

Ryan Boldi, Lijie Ding, Lee Spector, et al. (2024). *Pareto-Optimal Learning from Preferences with Hidden Context*. arXiv.org.

Yuhao Du, Zhuo Li, Pengyu Cheng, et al. (2025). *Simplify RLHF as Reward-Weighted SFT: A Variational Method*. arXiv.org.

Yuhang Lai, Siyuan Wang, Shujun Liu, et al. (2024). *ALaRM: Align Language Models via Hierarchical Rewards Modeling*. Annual Meeting of the Association for Computational Linguistics.

Han Zhang, Lin Gui, Yu Lei, et al. (2024). *COPR: Continual Human Preference Learning via Optimal Policy Regularization*. arXiv.org.

Somanshu Singla, Zhen Wang, Tianyang Liu, et al. (2024). *Dynamic Rewarding with Prompt Optimization Enables Tuning-free Self-Alignment of Language Models*. Conference on Empirical Methods in Natural Language Processing.

Ruichen Shao, Bei Li, Gangao Liu, et al. (2025). *Earlier Tokens Contribute More: Learning Direct Preference Optimization From Temporal Decay Perspective*. International Conference on Learning Representations.

Hongyu Yang, Liyang He, Min Hou, et al. (2024). *Aligning LLMs through Multi-perspective User Preference Ranking-based Feedback for Programming Question Answering*. arXiv.org.

Yekun Chai, Haoran Sun, Huang Fang, et al. (2024). *MA-RLHF: Reinforcement Learning from Human Feedback with Macro Actions*. International Conference on Learning Representations.

Avinash Anand, Kritarth Prasad, Chhavi Kirtani, et al. (2024). *Enhancing LLMs for Physics Problem-Solving using Reinforcement Learning with Human-AI Feedback*. arXiv.org.

Hao Sun, Yunyi Shen, Jean-Francois Ton, et al. (2025). *Reusing Embeddings: Reproducible Reward Model Research in Large Language Model Alignment without GPUs*. arXiv.org.

Zhilun Zhou, Yuming Lin, and Yong Li (2024). *Large language model empowered participatory urban planning*. arXiv.org.

Kuang-Ming Chen, and Hung-yi Lee (2024). *InstructionCP: A fast approach to transfer Large Language Models into target language*. arXiv.org.

Yang Chen, Yufan Shen, Wenxuan Huang, et al. (2025). *Learning Only with Images: Visual Reinforcement Learning with Reasoning, Rendering, and Visual Feedback*. arXiv.org.

Ning Wang, Bingkun Yao, Jie Zhou, et al. (2025). *Insights from Verification: Training a Verilog Generation LLM with Reinforcement Learning with Testbench Feedback*. arXiv.org.

D. Tiapkin, Daniele Calandriello, Johan Ferret, et al. (2025). *On Teacher Hacking in Language Model Distillation*. arXiv.org.

Longxi Gao, Li Zhang, and Mengwei Xu (2025). *UIShift: Enhancing VLM-based GUI Agents through Self-supervised Reinforcement Learning*. arXiv.org.

Graziano A. Manduzio, F. Galatolo, M. G. Cimino, et al. (2024). *Improving Small-Scale Large Language Models Function Calling for Reasoning Tasks*. arXiv.org.

Eduardo Pignatelli, Johan Ferret, Tim Rockaschel, et al. (2024). *Assessing the Zero-Shot Capabilities of LLMs for Action Evaluation in RL*. arXiv.org.

Qining Zhang, Honghao Wei, and Lei Ying (2024). *Reinforcement Learning from Human Feedback without Reward Inference: Model-Free Algorithm and Instance-Dependent Analysis*. RLJ.

Yu Zhu, Chuxiong Sun, Wenfei Yang, et al. (2024). *Proxy-RLHF: Decoupling Generation and Alignment in Large Language Model with Proxy*. arXiv.org.

S. Srivastava, and Vaneet Aggarwal (2025). *A Technical Survey of Reinforcement Learning Techniques for Large Language Models*. arXiv.org.

Debmalya Mandal, Paulius Sasnauskas, and Goran Radanovic (2025). *Distributionally Robust Reinforcement Learning with Human Feedback*. arXiv.org.

Evan Chen, Run-Jun Zhan, Yan-Bai Lin, et al. (2025). *More Women, Same Stereotypes: Unpacking the Gender Bias Paradox in Large Language Models*. Unpublished manuscript.

Heyang Zhao, Chen Ye, Wei Xiong, et al. (2025). *Logarithmic Regret for Online KL-Regularized Reinforcement Learning*. arXiv.org.

Johan S Daniel, and Anand Pal (2024). *Impact of Non-Standard Unicode Characters on Security and Comprehension in Large Language Models*. arXiv.org.

Jiahui Li, Lin Li, Tai-wei Chang, et al. (2024). *RED: Unleashing Token-Level Rewards from Holistic Feedback via Reward Redistribution*. Unpublished manuscript.

_1_historical_trajectory__and__early.promise_of_rl_in_nlp1.13Historical Trajectory and Early Promise of RL in NLPsubsection.1.1 _2_contemporary_motivations:_aligning_Motivations: Aligning Large Language Models with Human Intentssubsection.1.2 _3_scope_and_Structure_of_the_Reviewsubsection.1.3 _concepts:_bridging_rl__and__language211Found Concepts: Bridging RL and Languagesection.2 _1_core_reinforcement_learning_algorithms_for_seq Reinforcement Learning Algorithms for Sequential Decision Makingsubsection.2.1 _2_structured_prediction_as_a_search_problem2.213Structured Prediction as a Search Problemsubsection.2.2 _3_optimizing_non-differentiable_metrics_in_nlp2.315Optimizing Non-Differentiable Metrics in NLPsubsection.2.3 _reinforcement_learning_for_direct_language_gen Reinforcement Learning for Direct Language Generationsection.3 _1_drl_for_dialogue_generation3.11for_Dialogue_Generationsubsection.3.1 _2_drl_for_abstractive_summarization__and__machine_for_Abstractive_Summarization_and_Machine_Translationsubsection.3.2 _3_actor-critic_methods__and__exposure_bias_mitigation3.32Actor-Critic Methods and Exposure Bias Mitigationsubsection.3.3 _rise_of_reinforcement_learning_from_human_feedback Rise of Reinforcement Learning from Human Feedback (RLHF)section.4 _1_core_rlhf_paradigm:_reward RLHF Paradigm: Reward Modeling and Policy Optimizationsubsection.4.1 _2_proximal_policy_optimization_(ppo)_for_llm_alignment4.227Proximal Policy Optimization (PPO) for LLM Alignmentsubsection.4.2 _3_direct_preference_optimization_(dpo)__and__reward-model-free_alignment4.330Direct Preference Optimization (DPO) and Reward-Model-Free Alignmentsubsection.4.3 _4_theoretical_foundations_of_kl-constrained_preference_learning Foundations of KL-Constrained Preference Learningsubsection.4.4 _feed-back_mechanisms__and__robust_alignment536Advanced Feedback Mechanisms and Robust Alignmentsection.5 _1_reinforcement_learning_from_ai_feedback_(rlaif)5.136Reinforcement Learning from AI Feedback (RLAIF)subsection.5.1 _2_mitigating_reward_model_imperfections Reward Model Imperfections and Overoptimizationsubsection.5.2 _3_interpretable__and__multi-objective_reward_modeling5.341Interpretable and Multi-Objective Reward Modeling-subsection.5.3 _4_implicit_human_signals__and__distributionally_robust_alignment5.444Implicit Human Signals and Distributionally Robust Alignmentsubsection.5.4 _for_enhanced_llm_capabilities for Enhanced LLM Capabilities and Specialized Applicationssection.6 _1_enhancing_llm_reasoning_v